

Linear Models in R (M1–MIDO)

Lab Session 1 – Solutions

Henri PANJO

Table of contents

Dataset Overview: <i>data_pokemon.csv</i>	3
Setup	4
Question 1. Loading dataset	5
Solutions	5
Question 2. Summary statistics	7
Solutions	7
Question 3. Histogram and Scatter plots	10
Solutions	10
Question 4. Correlation matrix	13
Solutions	13
Question 5. Simple Linear Regression (SLR)	15
Solutions	15
Question 6. Fitted values, residuals	20
Solutions	20
Question 7. Residual diagnostics	22
Note: Standardized vs Studentized residuals	22
Solutions	24
Question 8. Prediction & Intervals	36
Solutions	36
Session Info	39

Dataset Overview: `data_pokemon.csv`

This dataset is adapted from a popular Kaggle Pokémon dataset.

Even if you are not familiar with Pokémon, the data is straightforward:

it combines numeric statistics with categorical attributes, making it well-suited for applying Ordinary Least Squares (OLS) in R.

What it contains

- Unique identifiers and names for each Pokémon
- Battle statistics (health, attack, defense, special attack, special defense, speed)
- Categorical features (primary/secondary type, generation, legendary flag)

Fields (Codebook)

- `id`: Unique Pokémon ID
- `name`: Pokémon name
- `type_1`: Primary type (e.g., Water, Fire)
- `type_2`: Secondary type (optional)
- `hp`: Hit points (overall health)
- `attack`: Physical attack strength (we will use this as y in most regressions)
- `defense`: Physical defense strength
- `sp_attack`: Special (non-physical) attack strength
- `sp_defense`: Special defense strength
- `speed`: Speed / turn order
- `generation`: Game generation label
- `legendary`: Indicator for legendary status (TRUE/FALSE)

Note on notation

- We treat `attack` as the outcome variable Y .
- Predictor variables (e.g., `defense`, `speed`) will be denoted as x_1, x_2, \dots .
- Factors like `type_1` or `legendary` will be included as categorical predictors.

Setup

To keep numbers readable and reproducible, we set display options:

```
options(scipen = 999, digits = 5)
```

We also load the packages used during this session.

i Note

Don't worry if you don't know them all — we'll introduce functions as we need them. Some provide regression tools, others are for data visualization or diagnostics.

```
library(broom)
library(performance)
library(parameters)
library(datawizard)
library(see)
library(effectsize)
library(insight)
library(correlation)
library(modelbased)
library(glue)
library(scales)
library(GGally)
library(ggpubr)
library(car)
library(lmtest)
library(rstatix)
library(matrixTests)
library(ggfortify)
library(qqplotr)
library(collapse)
library(tidyverse)
```

Question 1. Loading dataset

Import the `data_pokemon.csv` file with `read_csv()`. Save the data in an object called `pok`.

- Quickly examine the data using `glimpse()` from `{dplyr}`
 - Display the first 10 rows of `pok` using `head()` or `slice()`.

Solutions

- Loading data_pokemon.csv

```
pok <- read_csv("data_pokemon.csv", show_col_types = FALSE)
```

- `glimpse()` on pok

glimpse(pok)

- `head()` and `slice()` on `pok`

`head(pok, n = 10)`

```
# A tibble: 10 x 12
  id name   type_1 type_2   hp attack defense sp_attack sp_def speed generation legendary
  <dbl> <chr>   <chr>   <chr> <dbl> <dbl>   <dbl>    <dbl> <dbl> <dbl>   <dbl> <chr>
1 1 Bulbasaur Grass Poison 45 49 49 65 65 45 1 No
2 2 Ivysaur   Grass Poison 60 62 63 80 80 60 1 No
3 3 Venusaur  Grass Poison 80 82 83 100 100 80 1 No
4 4 Mega Ven~ Grass Poison 80 100 123 122 120 80 1 No
5 5 Charmand~ Fire  None 39 52 43 60 50 65 1 No
6 6 Charmele~ Fire  None 58 64 58 80 65 80 1 No
7 7 Charizard  Fire Flying 78 84 78 109 85 100 1 No
8 8 Mega Cha~ Fire Dragon 78 130 111 130 85 100 1 No
9 9 Mega Cha~ Fire Flying 78 104 78 159 115 100 1 No
10 10 Squirtle Water None 44 48 65 50 64 43 1 No
```

`slice(pok, 1:10)`

```
# A tibble: 10 x 12
  id name   type_1 type_2   hp attack defense sp_attack sp_def speed generation legendary
  <dbl> <chr>   <chr>   <chr> <dbl> <dbl>   <dbl>    <dbl> <dbl> <dbl>   <dbl> <chr>
1 1 Bulbasaur Grass Poison 45 49 49 65 65 45 1 No
2 2 Ivysaur   Grass Poison 60 62 63 80 80 60 1 No
3 3 Venusaur  Grass Poison 80 82 83 100 100 80 1 No
4 4 Mega Ven~ Grass Poison 80 100 123 122 120 80 1 No
5 5 Charmand~ Fire  None 39 52 43 60 50 65 1 No
6 6 Charmele~ Fire  None 58 64 58 80 65 80 1 No
7 7 Charizard  Fire Flying 78 84 78 109 85 100 1 No
8 8 Mega Cha~ Fire Dragon 78 130 111 130 85 100 1 No
9 9 Mega Cha~ Fire Flying 78 104 78 159 115 100 1 No
10 10 Squirtle Water None 44 48 65 50 64 43 1 No
```

Question 2. Summary statistics

For the variables attack, speed, defense, hp, compute summary statistics: number of missing values, number of distinct values, mean, median, and standard deviation.

Hint: `summary()`, `descr()`, `describe_distribution()`, `get_summary_stats()`, `summarise()`, `mean()`, `sd()`, `median()`, `n_distinct()`, `is.na()`

Solutions

- `summary()` on selected variables attack, speed, defense, hp

```
select(pok, attack, speed, defense, hp) |>  
  summary()
```

attack	speed	defense	hp
Min. : 5	Min. : 5.0	Min. : 5.0	Min. : 1.0
1st Qu.: 55	1st Qu.: 45.0	1st Qu.: 50.0	1st Qu.: 50.0
Median : 75	Median : 65.0	Median : 70.0	Median : 65.0
Mean : 79	Mean : 68.3	Mean : 73.8	Mean : 69.3
3rd Qu.: 100	3rd Qu.: 90.0	3rd Qu.: 90.0	3rd Qu.: 80.0
Max. : 190	Max. : 180.0	Max. : 230.0	Max. : 255.0

- `descr()` (`{collapse}`)

```
select(pok, attack, speed, defense, hp) |>  
  descr(Ndistinct = TRUE, Qprobs = c(0.25, 0.5, 0.75)) |>  
  as_tibble()
```

	# A tibble: 4 x 13	Variable	Class	N	Ndist	Mean	SD	Min	Max	Skew	Kurt	`25%`	`50%`	`75%`
		<chr>	<chr>	<dbl>										
1	attack	numeric	800	111	79.0	32.5	5	190	0.551	3.16	55	75	100	
2	speed	numeric	800	108	68.3	29.1	5	180	0.357	2.76	45	65	90	
3	defense	numeric	800	103	73.8	31.2	5	230	1.15	5.70	50	70	90	
4	hp	numeric	800	94	69.3	25.5	1	255	1.57	10.2	50	65	80	

- `describe_distribution()` (`{datawizard}`)

```
select(pok, attack, speed, defense, hp) |>
  describe_distribution(centrality = c("mean", "median"), quartiles = TRUE) |>
  as_tibble()
```

	Variable	Median	MAD	Mean	SD	IQR	Min	Max	Q1	Q3	Skewness	Kurtosis	n
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<int>
1	attack	75	29.7	79.0	32.5	45	5	190	55	100	0.552	0.170	800
2	speed	65	31.1	68.3	29.1	45	5	180	45	90	0.358	-0.236	800
3	defense	70	29.7	73.8	31.2	40	5	230	50	90	1.16	2.73	800
4	hp	65	22.2	69.3	25.5	30	1	255	50	80	1.57	7.23	800

- `get_summary_stats()` (`{rstatix}`)

```
select(pok, attack, speed, defense, hp) |>
  get_summary_stats(show = c("n", "mean", "sd", "median", "q1", "q3"))
```

	variable	n	mean	sd	median	q1	q3
	<fct>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	attack	800	79.0	32.5	75	55	100
2	speed	800	68.3	29.1	65	45	90
3	defense	800	73.8	31.2	70	50	90
4	hp	800	69.3	25.5	65	50	80

- `summarise()` (`{dplyr}`). More complicated

- First we create a list of functions

```
myfunctions <- list(
  n = length, nmiss = \((x)\) sum(is.na(x)), ndistinct = n_distinct,
  mean = mean, sd = sd, median = median
)
```

- We use `summarise()` with `across()` and `myfunctions`. Then we apply `pivot_longer()` (`{tidyverse}`)

```
pok |>
  summarise(across(c("attack", "speed", "defense", "hp"), myfunctions)) |>
  pivot_longer(
    cols = everything(),
    names_to = c("Variable", ".value"),
    names_sep = "_"
  )
```

```
# A tibble: 4 x 7
  Variable     n nmiss ndistinct   mean     sd median
  <chr>     <int> <int>     <int> <dbl> <dbl> <dbl>
1 attack      800     0         111  79.0  32.5    75
2 speed       800     0         108  68.3  29.1    65
3 defense     800     0         103  73.8  31.2    70
4 hp          800     0         94   69.3  25.5    65
```


Question 3. Histogram and Scatter plots

1. Plot histogram of attack and speed. Hint: `geom_histogram()`
2. Create scatter plots of attack against each numeric predictor speed, defense, hp. Hint: `geom_point()`, with `geom_smooth(method = "lm")`

Solutions


- `ggplot()` and `geom_histogram()` on the vector attack from the dataframe pok

```
ggplot(pok, aes(x = speed)) +  
  geom_histogram(bins = 10, color = "black", fill = "dodgerblue") +  
  scale_y_continuous(expand = expansion(c(0, 0.05))) +  
  labs(x = "Speed", y = "Count") +  
  theme_bw(base_size = 14) +  
  labs_pubr()
```




- Scatter plot of attack against speed

```
pok |>  
  ggplot(aes(x = speed, y = attack)) +  
  geom_point(size = 2, shape = 21, fill = "dodgerblue", color = "black") +  
  labs(x = "Speed (x)", y = "Attack (y)") +  
  theme_bw(base_size = 14) +  
  labs_pubr()
```



- Scatter plots of attack against each numeric predictor speed, defense, hp

```
select(pok, attack, speed, defense, hp) |>
pivot_longer(cols = c(speed, defense, hp)) |>
ggplot(aes(x = value, y = attack)) +
facet_wrap(vars(name), ncol = 1) +
geom_point(size = 2, shape = 21, fill = "red", color = "black", alpha = 1) +
labs(x = "Value (x)", y = "Attack (y)") +
theme_bw(base_size = 14) +
labs_pubr()
```



Question 4. Correlation matrix

Compute and interpret the correlation matrix of the predictors speed, defense, hp.

Hint: `cor()`, `correlation()`

Solutions

- With `cor()`

```
select(pok, speed, defense, hp) |>  
  cor(method = "pearson")
```

```
      speed  defense      hp  
speed  1.000000  0.015227  0.17595  
defense 0.015227  1.000000  0.23962  
hp      0.175952  0.239622  1.00000
```

- With `correlation()` from `{correlation}`

```
select(pok, speed, defense, hp) |> correlation(method = "pearson")
```

```
# Correlation Matrix (pearson-method)  
  
Parameter1 | Parameter2 |     r |      95% CI | t(798) |      p  
-----  
speed     | defense    | 0.02 | [-0.05, 0.08] |  0.43 | 0.667  
speed     | hp         | 0.18 | [ 0.11, 0.24] |  5.05 | < .001***  
defense   | hp         | 0.24 | [ 0.17, 0.30] |  6.97 | < .001***  
  
p-value adjustment method: Holm (1979)  
Observations: 800
```


```
select(pok, speed, defense, hp) |>  
  correlation(method = "pearson") |>  
  summary(redundant = TRUE)
```

```
# Correlation Matrix (pearson-method)  
  
Parameter |     speed | defense |      hp  
-----  
speed     |           | 0.02 | 0.18***  
defense   | 0.02 |       | 0.24***  
hp       | 0.18*** | 0.24*** |  
  
p-value adjustment method: Holm (1979)
```

- We can have a plot

```
select(pok, speed, defense, hp) |>
  correlation(method = "pearson") |>
  summary(redundant = TRUE) |>
  plot() +
  theme_bw(base_size = 14) +
  theme(legend.position = "bottom")
```

Correlation Matrix



Correlation 0.00 0.25 0.50 0.75 1.00

Question 5. Simple Linear Regression (SLR)

1. For each predictors speed, defense, hp, fit a SLR to explain the variable attack:

$$y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

$$\text{attack} = \beta_0 + \beta_1 \text{speed} + \varepsilon$$

$$\text{attack} = \beta_0 + \beta_1 \text{defense} + \varepsilon$$

$$\text{attack} = \beta_0 + \beta_1 \text{hp} + \varepsilon$$

Save the models in 3 objects: `slr_speed`, `slr_defense`, `slr_hp`. Interpret intercept and slope.

Hint: `lm()`

2. Interpret the results of the three hypothesis tests $H_0 : \beta_j = 0$ versus $H_1 : \beta_j \neq 0$?

Hint: `summary()` or `coeftest()` on `slr_speed`, `slr_defense`, `slr_hp`

3. Compute 95% CIs for coefficients.

Hint: `confint()`, `tidy()`, `model_parameters()`

Solutions

- We use `lm()` to fit the simple regressions models

```
slr_speed <- lm(formula = attack ~ speed, data = pok)
slr_speed
```

```
Call:
lm(formula = attack ~ speed, data = pok)
```

```
Coefficients:
(Intercept)      speed
 49.928        0.426
```

$$\widehat{\text{attack}} = 49.928 + 0.426 \times \text{speed}$$

```
slr_defense <- lm(attack ~ defense, data = pok)
slr_defense
```

```
Call:
lm(formula = attack ~ defense, data = pok)
```

```
Coefficients:
(Intercept)      defense
        45.284          0.457
```

$$\widehat{\text{attack}} = 45.284 + 0.457 \times \text{defense}$$

```
slr_hp <- lm(attack ~ hp, data = pok)
slr_hp
```

```
Call:
lm(formula = attack ~ hp, data = pok)
```

```
Coefficients:
(Intercept)      hp
        41.816          0.537
```

$$\widehat{\text{attack}} = 41.816 + 0.537 \times \text{hp}$$

- What kind of object does `lm()` create ?

```
class(slr_hp)
```

```
[1] "lm"
```

```
typeof(slr_hp)
```

```
[1] "list"
```

```
names(slr_hp)
```

```
[1] "coefficients"   "residuals"       "effects"         "rank"
[7] "qr"              "df.residual"    "xlevels"        "call"
                                         "fitted.values" "terms"          "model"
```

```
sum_lm_speed <- summary(slr_speed)
sum_lm_speed
```

```
Call:
lm(formula = attack ~ speed, data = pok)

Residuals:
    Min      1Q  Median      3Q     Max 
-73.99 -21.55   -3.79   18.07 103.14 

Coefficients:
            Estimate Std. Error t value     Pr(>|t|)    
(Intercept) 49.9285    2.7120   18.4 <0.000000000000002 *** 
speed       0.4258    0.0366   11.6 <0.000000000000002 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 

Residual standard error: 30 on 798 degrees of freedom
Multiple R-squared:  0.145, Adjusted R-squared:  0.144 
F-statistic: 136 on 1 and 798 DF,  p-value: <0.000000000000002
```

- What kind of object does `summary()` create?

```
class(sum_lm_speed)
```

```
[1] "summary.lm"
```

```
typeof(sum_lm_speed)
```

```
[1] "list"
```

```
names(sum_lm_speed)
```

```
[1] "call"           "terms"        "residuals"      "coefficients" "aliased"      "sigma"      
[7] "df"             "r.squared"    "adj.r.squared" "fstatistic"    "cov.unscaled"
```

- Test w of $\beta = 0$ with `coeftest()` from `{lmtest()}`

```
coeftest(slr_speed)
```

```
t test of coefficients:

      Estimate Std. Error t value      Pr(>|t|)    
(Intercept) 49.9285     2.7120    18.4 <0.0000000000000002 *** 
speed       0.4258     0.0366    11.6 <0.0000000000000002 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coeftest(slr_defense)
```

```
t test of coefficients:

      Estimate Std. Error t value      Pr(>|t|)    
(Intercept) 45.2842     2.6538    17.1 <0.0000000000000002 *** 
defense      0.4566     0.0331    13.8 <0.0000000000000002 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
coeftest(slr_hp)
```

```
t test of coefficients:

      Estimate Std. Error t value      Pr(>|t|)    
(Intercept) 41.8163     3.0104    13.9 <0.0000000000000002 *** 
hp          0.5369     0.0408    13.2 <0.0000000000000002 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- 95% CIs for coefficients of slr_speed with `confint()`

```
confint(slr_speed, level = 0.95)
```

	2.5 %	97.5 %
(Intercept)	44.60493	55.25204
speed	0.35405	0.49755

- 95% CIs for coefficients of slr_defense with `tidy()` from `{broom}`

```
tidy(slr_defense, conf.int = TRUE, conf.level = 0.95)
```

# A tibble: 2 x 7	term	estimate	std.error	statistic	p.value	conf.low	conf.high
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	45.3	2.65	17.1	6.82e-56	40.1	50.5
2	defense	0.457	0.0331	13.8	5.86e-39	0.392	0.522

- 95% CIs for coefficients of slr_hp with `model_parameters()` from `{parameters}`

```
model_parameters(slr_hp, ci = 0.95, ci_method = "residual", digits = 3)
```

Parameter	Coefficient	SE	95% CI	t(798)	p
(Intercept)	41.816	3.010	[35.907, 47.726]	13.891	< .001
hp	0.537	0.041	[0.457, 0.617]	13.164	< .001

Question 6. Fitted values, residuals

1. In the pok database, create the following variables

- `yhat_speed`, which represents the fitted values of the `slr_speed` model: $\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_i$
- `res_speed`, which represents the residuals of the `slr_speed` model: $\hat{\varepsilon}_i = y_i - \hat{y}_i$

Hint: retrieve $\hat{\beta}_0$ and $\hat{\beta}_1$ with `coef()`

Display the first 10 rows of `pok` with the fitted values and residuals added.

2. Compute the fitted values and residuals using the functions `predict()` (or `fitted`) and `residuals()` (or `resid()`) on `slr_speed`. Also try the handy function `augment()` on `slr_speed`

Solutions

- We retrieve $\hat{\beta}_0$ and $\hat{\beta}_1$ with `coef()`

```
betas_speed <- coef(slr_speed)
betas_speed
```

```
(Intercept)      speed
49.9285        0.4258
```

- We create `yhat_speed` and `res_speed` in `pok` with `mutate()` from `{dplyr}`

```
pok <- pok |>
  mutate(yhat_speed = betas_speed[1] + betas_speed[2] * speed) |>
  mutate(res_speed = attack - yhat_speed)

select(pok, id, name, attack, speed, yhat_speed, res_speed) |> head(10)
```

```
# A tibble: 10 x 6
  id   name      attack  speed  yhat_speed  res_speed
  <dbl> <chr>     <dbl>  <dbl>     <dbl>      <dbl>
1    1 Bulbasaur     49     45     69.1     -20.1
2    2 Ivysaur       62     60     75.5     -13.5
3    3 Venusaur      82     80     84.0     -1.99
4    4 Mega Venusaur 100     80     84.0      16.0
5    5 Charmander     52     65     77.6     -25.6
6    6 Charmeleon     64     80     84.0     -20.0
7    7 Charizard      84    100     92.5     -8.51
8    8 Mega Charizard X 130    100     92.5      37.5
9    9 Mega Charizard Y 104    100     92.5      11.5
10   10 Squirtle      48     43     68.2     -20.2
```

- With `fitted()` and `residuals()`. 10 first values

```
fitted(slr_speed)[1:10]
```

1	2	3	4	5	6	7	8	9	10
69.090	75.477	83.993	83.993	77.606	83.993	92.509	92.509	92.509	68.238

```
residuals(slr_speed)[1:10]
```

1	2	3	4	5	6	7	8	9	10
-20.0896	-13.4767	-1.9927	16.0073	-25.6057	-19.9927	-8.5088	37.4912	11.4912	-20.2380

- With `augment()` from `{broom}`

```
augment_speed <- augment(slr_speed)
head(augment_speed, 10)
```

	# A tibble: 10 x 8	attack	speed	.fitted	.resid	.hat	.sigma	.cooksdi	.std.resid
		<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	49	45	69.1	-20.1	0.00205	30.0	0.000461	-0.670	
2	62	60	75.5	-13.5	0.00135	30.0	0.000137	-0.449	
3	82	80	84.0	-1.99	0.00145	30.0	0.00000321	-0.0664	
4	100	80	84.0	16.0	0.00145	30.0	0.000207	0.534	
5	52	65	77.6	-25.6	0.00127	30.0	0.000462	-0.853	
6	64	80	84.0	-20.0	0.00145	30.0	0.000323	-0.666	
7	84	100	92.5	-8.51	0.00274	30.0	0.000111	-0.284	
8	130	100	92.5	37.5	0.00274	30.0	0.00215	1.25	
9	104	100	92.5	11.5	0.00274	30.0	0.000202	0.383	
10	48	43	68.2	-20.2	0.00220	30.0	0.000501	-0.675	

Question 7. Residual diagnostics

Note: Standardized vs Studentized residuals

$$\hat{\varepsilon} = \mathbf{y} - \hat{\mathbf{y}} = \mathbf{y} - \mathbb{X}\hat{\beta} = \mathbf{y} - P_{\mathbb{X}}\mathbf{y} = (\mathbf{I}_n - P_{\mathbb{X}})\mathbf{y} = (\mathbf{I}_n - P_{\mathbb{X}})\boldsymbol{\varepsilon}$$

- Let denote by h_{ij} the element of the projector $P_{\mathbb{X}} = H_{\mathbb{X}}$ such that $P_{\mathbb{X}} = H_{\mathbb{X}} = [h_{ij}]$
- The diagonal elements $h_{ii} \in [0, 1]$ are called the *leverages*
- If $h_{ii} > 2p/n$ (sometimes $h_{ii} > 3p/n$), then the observation i is consider an *outlier*
- We have $\text{Cov}(\boldsymbol{\varepsilon}) = \sigma^2 \mathbf{I}_n$ but $\text{Cov}(\hat{\varepsilon}) = \sigma^2 (\mathbf{I}_n - H_{\mathbb{X}})$
- The residuals are not independant, however, in many cases, especially if n is large, the h_{ii} 's tend to be small.
The impact of this is usually small and diagnostics can reasonably be applied to the residuals in order to check the assumptions on the error but we can also modify the residuals to adjust for this effect.
- Standardized residuals (from `rstandard()`)
Raw residuals are rescaled by their estimated standard deviation, taking into account leverage.

$$\hat{r}_i = \frac{\hat{\varepsilon}_i}{\hat{\sigma} \sqrt{1 - h_{ii}}}$$

where $\hat{\varepsilon}_i$ is the raw residual and h_{ii} is the leverage of observation i .
These make residuals roughly comparable across observations.

- Studentized residuals (from `rstudent()`)
Go one step further: each residual is scaled using a variance estimate that excludes the i -th observation.
This gives more accurate standard errors and makes large outliers easier to detect.

$$t_i^* = \frac{\hat{\varepsilon}_i}{\hat{\sigma}_{(-i)} \sqrt{1 - h_{ii}}}$$

where $\hat{\sigma}_{(-i)}$ is the error standard deviation estimated without observation i .

In practice:


- Use standardized residuals for quick checks.
- Use studentized residuals when formally testing for outliers or influential observations.

1. For the `s1r_speed` model, with the help of `autoplot()`, plot residuals vs fitted values and $\sqrt{|\text{Standardized residuals}|}$ vs fitted values
2. Plot *studentized residuals* vs fitted: Hint: `fitted()`, `rstudent()`
3. Plot residuals vs `speed` and vs `defense`
4. Plot residuals in the order of observation (to detect time dependence). Hint: `augment()`
5. Plot an histogram of the standardized residuals.
6. Plot a normal Q-Q plot of the standardized residuals. Hint: `qqnorm()`, `qqline()`, `augment()`, `stat_qq()`, `stat_qq_line()`
7. Performs the Breusch–Pagan test for heteroskedasticity: Hint: `bptest()`, `ncvTest()`
8. Perform the Durbin–Watson test on the residuals: Hint: `dwttest()`, `durbinWatsonTest()`
9. Perform test for normality on standardized residuals. Hint: `shapiro.test()`, `shapiro_test()`, `col_jarquebera()`

Solutions

- Residuals vs fitted values with `autoplot()` (need to load `{ggfortify}`)


```
autoplot(slr_speed, which = 1, ncol = 1, colour = "black", shape = 21, size = 2) +  
  theme_bw(base_size = 14) +  
  labs_pubr()
```



- $\sqrt{|\text{Standardized residuals}|}$ vs fitted

```
autoplot(slr_speed, which = 3, ncol = 1, colour = "black", shape = 21, size = 2) +
  theme_bw(base_size = 14) +
  labs_pubr()
```

Scale-Location




- Homoscedasticity is satisfied

- We create a dataset with fitted values and *studentized residuals* (`rstudent()`)

```
tibble(
  fitted = fitted(slr_speed), rstudent = rstudent(slr_speed)
) |>
  ggplot(aes(x = fitted, y = rstudent)) +
  geom_hline(yintercept = 0, linetype = 2) +
  geom_point(shape = 21, size = 2) +
  geom_smooth(method = "loess", se = FALSE) +
  labs(x = "Fitted values", y = "Internally studentized residuals") +
  labs(title = "Studentized residuals vs Fitted values") +
  theme_bw(base_size = 14) +
  labs_pubr()
```


Studentized residuals vs Fitted values



- Linearity and homoscedasticity are satisfied

- Residuals vs speed


```
pok |>
  ggplot(aes(x = speed, y = res_speed)) +
  geom_hline(yintercept = 0, linetype = 2) +
  geom_point(shape = 21, size = 2) +
  geom_smooth(method = "loess", se = FALSE) +
  labs(x = "Speed", y = "Residuals") +
  labs(title = "Residuals vs Speed") +
  theme_bw(base_size = 14) +
  labs_pubr()
```



- Linearity and homoscedasticity are satisfied

- Speed Residuals vs defense

```
pok |>
  ggplot(aes(x = defense, y = res_speed)) +
  geom_hline(yintercept = 0, linetype = 2) +
  geom_point(shape = 21, size = 2) +
  geom_smooth(method = "loess", se = FALSE) +
  labs(x = "Defense", y = "speed Residuals") +
  labs(title = "Speed Residuals vs Defense") +
  theme_bw(base_size = 14) +
  labs_pubr()
```




- We should probably add defense to the model

- Residuals vs the observation order


```
pok |>
  ggplot(aes(x = id, y = res_speed)) +
  geom_hline(yintercept = 0, linetype = 2) +
  geom_point(shape = 21, size = 2) +
  geom_smooth(method = "loess", se = FALSE) +
  labs(x = "Observation order", y = "Residuals") +
  labs(title = "Residuals vs Observation order") +
  theme_bw(base_size = 14) +
  labs_pubr()
```

Residuals vs Observation order




- Histogram of the standardized residuals

```
augment_speed |>
  ggplot(aes(x = .std.resid)) +
  geom_histogram(bins = 12, color = "black", fill = "dodgerblue") +
  scale_y_continuous(expand = expansion(c(0, 0.05))) +
  labs(x = "Standardized residuals", y = "Count") +
  theme_bw(base_size = 14) +
  labs_pubr()
```




- Density of the standardized residuals

```
ggdensity(augment_speed, x = ".std.resid", fill = "dodgerblue") +  
  scale_x_continuous(limits = c(-4, 4)) +  
  stat_overlay_normal_density(color = "red", linetype = 1, linewidth = 1) +  
  scale_y_continuous(expand = expansion(c(0, 0.05))) +  
  labs(x = "Standardized residuals", y = "Density") +  
  theme_bw(base_size = 14) +  
  labs_pubr()
```



- Normal Q-Q plot of the standardized residuals

```
augment_speed |>
  ggplot(aes(sample = .std.resid)) +
  stat_qq_band(alpha = 0.2, fill = "blue") + # du package {qqplotr}
  stat_qq_line(color = "red") + # version du package {qqplotr}
  stat_qq_point(size = 0.5) +
  labs(y = "Sample quantile", x = "Theoretical quantile") +
  theme_bw(base_size = 14) +
  labs_pubr()
```



- Breusch–Pagan test for heteroskedasticity with `bptest()` from `{lmtest}` or `ncvTest()` from `{car}`

```
bptest(slr_speed, studentize = FALSE)
```

Breusch-Pagan test

```
data: slr_speed  
BP = 7.92, df = 1, p-value = 0.0049
```

```
ncvTest(slr_speed)
```

```
Non-constant Variance Score Test  
Variance formula: ~ fitted.values  
Chisquare = 7.9215, Df = 1, p = 0.00489
```

- Durbin–Watson test with `dwtest()` from `{lmtest}` or `durbinWatsonTest()` from `{car}`

```
dwtest(slr_speed)
```

Durbin-Watson test

```
data: slr_speed  
DW = 1.4, p-value <0.0000000000000002  
alternative hypothesis: true autocorrelation is greater than 0
```

```
durbinWatsonTest(slr_speed)
```

```
lag Autocorrelation D-W Statistic p-value  
1      0.29892      1.4003      0  
Alternative hypothesis: rho != 0
```

- Shapiro-Wilk Normality Test on standardized residuals

```
shapiro.test(augment_speed[["std.resid"]])
```

```
Shapiro-Wilk normality test  
data: augment_speed[["std.resid"]]  
W = 0.982, p-value = 0.000000025
```

```
augment_speed |>  
  shapiro_test(.std.resid)
```

```
# A tibble: 1 x 3  
  variable   statistic      p  
  <chr>       <dbl>     <dbl>  
1 .std.resid    0.982 0.000000245
```

- Jarque-Bera test on standardized residuals

```
select(augment_speed, .std.resid) |>  
  col_jarquebera()
```


```
obs skewness kurtosis df statistic      pvalue  
.std.resid 800  0.50195  3.0528  2     33.686 0.000000048432
```

- `check_model()` from `{performance}`

```
check_model(slr_speed, check = c("qq", "normality", "linearity", "homogeneity"), size_dot = 1)
```


Linearity

Reference line should be flat and horizontal




Homogeneity of Variance

Reference line should be flat and horizontal




Normality of Residuals

Dots should fall along the line



Normality of Residuals

Distribution should be close to the normal curve



Question 8. Prediction & Intervals

Consider the model `slr_speed`

1. Compute 95% Confidence Interval for the mean $\mathbb{E}(\text{attack})$ at speed = 30, 70, 110, 150.

Hint: `predict(..., interval = "confidence")`, `estimate_expectation()`

2. Compute 95% Prediction Interval for new responses of attack at speed = 20, 60, 100, 140.

Hint: `predict(..., interval = "prediction")`, `estimate_prediction()`

3. On the graph representing the scatter plot between attack and speed, overlay the regression line and its 95% confidence band.

Solutions

- 95% Confidence Interval for the mean with `predict()`

```
grid_speed <- tibble(speed = seq(from = 30, to = 150, by = 40))

predict(slr_speed, newdata = grid_speed, interval = "confidence") |>
  as_tibble() |>
  mutate(speed = seq(30, 150, 40), .before = 1)
```

```
# A tibble: 4 x 4
  speed    fit   lwr   upr
  <dbl>   <dbl> <dbl> <dbl>
1     30    62.7  59.3  66.1
2     70    79.7  77.6  81.8
3    110   96.8  93.1 100.
4    150  114.   108.  120.
```

- With `estimate_expectation()` from `{modelbased}`

```
estimate_expectation(slr_speed, by = "speed = seq(30, 150, 40)", ci = 0.95) |>
  as_tibble()
```

```
# A tibble: 4 x 5
  speed Predicted     SE CI_low CI_high
  <dbl>     <dbl> <dbl>  <dbl>   <dbl>
1     30      62.7  1.76   59.3   66.1
2     70      79.7  1.06   77.6   81.8
3    110     96.8  1.86   93.1   100.
4    150    114.   3.17   108.   120.
```

- 95% Prediction Interval for the new observations with `predict()`

```
predict(slr_speed, newdata = grid_speed, interval = "prediction") |>
  as_tibble() |>
  mutate(speed = seq(30, 150, 40), .before = 1)
```

```
# A tibble: 4 x 4
  speed   fit    lwr    upr
  <dbl> <dbl> <dbl> <dbl>
1     30  62.7  3.66 122.
2     70  79.7 20.8 139.
3    110  96.8 37.7 156.
4    150 114.  54.5 173.
```

- With `estimate_prediction()` from `{modelbased}`

```
estimate_prediction(slr_speed, by = "speed = seq(30, 150, 40)", ci = 0.95)
```

Model-based Predictions


speed	Predicted	SE	95% CI
<hr/>			
30	62.70	30.08	[3.66, 121.74]
70	79.73	30.04	[20.76, 138.71]
110	96.77	30.08	[37.72, 155.82]
150	113.80	30.19	[54.53, 173.06]

Variable predicted: attack

Predictors modulated: speed = seq(30, 150, 40)

- Regression line and the 95% confidence interval of the predicted values

```
predict(slr_speed, newdata = select(pok, speed), interval = "confidence") |>
  as_tibble() |>
  bind_cols(select(pok, attack, speed)) |>
  ggplot(aes(x = speed, y = attack)) +
  geom_ribbon(aes(ymin = lwr, ymax = upr), alpha = 0.5, fill = "pink") +
  geom_point(size = 2, shape = 21, fill = "dodgerblue", color = "black") +
  geom_line(aes(y = fit), color = "black", linewidth = 0.5) +
  labs(x = "Speed (x)", y = "Attack (y)") +
  theme_bw(base_size = 14) +
  labs_pubr()
```



Session Info

- R version 4.5.1 (2025-06-13 ucrt)
- Rstudio version 2025.9.1.401 (Cucumberleaf Sunflower)

Package	Version
broom	1.0.10
car	3.1-3
collapse	2.1.3
correlation	0.8.8
datawizard	1.3.0
effectsize	1.0.1
GGally	2.4.0
ggfortify	0.4.19
ggpubr	0.6.1
glue	1.8.0
insight	1.4.2
lmtest	0.9-40
matrixTests	0.2.3
modelbased	0.13.0
parameters	0.28.2
performance	0.15.2
qqplotr	0.0.7
rstatix	0.7.2
scales	1.4.0
see	0.12.0
tidyverse	2.0.0