

Linear Models in R (M1–MIDO)

Lab Session 2 – Solutions

Henri PANJO

Table of contents

| | |
|---|----|
| Dataset Overview: <i>data_pokemon.csv</i> | 3 |
| Setup | 4 |
| Question 1. Loading dataset | 5 |
| Solutions | 5 |
| Question 2. Data management, label variable | 6 |
| Solutions | 6 |
| Question 3. Summary statistics | 7 |
| Solutions | 7 |
| Question 4. Histogram, Scatter plots, Correlations | 9 |
| Solutions | 9 |
| Question 5. Multivariate linear gaussian regression model | 12 |
| Solutions | 12 |
| Question 6. Test of overall regression | 16 |
| Solutions | 17 |
| Question 7. Indices of model performance for regression | 21 |
| Solutions | 21 |
| Question 8. Joint hypothesis test | 23 |
| Solutions | 23 |
| Question 9. Prediction & Intervals | 25 |
| Solutions | 25 |
| Question 10. Residual diagnostics | 27 |
| Note: Standardized vs Studentized residuals | 27 |
| Solutions | 28 |
| Session Info | 36 |

Dataset Overview: `data_pokemon.csv`

This dataset is adapted from a popular Kaggle Pokéémon dataset.

Even if you are not familiar with Pokéémon, the data is straightforward:

it combines numeric statistics with categorical attributes, making it well-suited for applying Ordinary Least Squares (OLS) in R.

What it contains

- Unique identifiers and names for each Pokéémon
- Battle statistics (health, attack, defense, special attack, special defense, speed)
- Categorical features (primary/secondary type, generation, legendary flag)

Fields (Codebook)

- `id`: Unique Pokéémon ID
- `name`: Pokéémon name
- `type_1`: Primary type (e.g., Water, Fire)
- `type_2`: Secondary type (optional)
- `hp`: Hit points (overall health)
- `attack`: Physical attack strength (we will use this as y in most regressions)
- `defense`: Physical defense strength
- `sp_attack`: Special (non-physical) attack strength
- `sp_defense`: Special defense strength
- `speed`: Speed / turn order
- `generation`: Game generation label
- `legendary`: Indicator for legendary status (TRUE/FALSE)

Note on notation

- We treat `attack` as the outcome variable Y .
- Predictor variables (e.g., `defense`, `speed`) will be denoted as x_1, x_2, \dots .
- Factors like `type_1` or `legendary` will be included as categorical predictors.

Setup

To keep numbers readable and reproducible, we set display options:

```
options(scipen = 999, digits = 5)
```

We also load the packages used during this session.

⚠ Warning

Don't worry if you don't know them all — we'll introduce functions as we need them. Some provide regression tools, others are for data visualization or diagnostics.

```
library(broom)
library(performance)
library(parameters)
library(datawizard)
library(see)
library(effectsize)
library(insight)
library(correlation)
library(modelbased)
library(glue)
library(scales)
library(GGally)
library(ggpubr)
library(car)
library(lmtest)
library(rstatix)
library(matrixTests)
library(ggfortify)
library(qqplotr)
library(patchwork)
library(gtsummary)
library(kableExtra)
library(collapse)
library(tidyverse)

source("helper_functions.R")
```

Question 1. Loading dataset

Import the dataset `data_pokemon.csv` with `read_csv()` and save it in an object called `pok`.

Using `select()`, keep only the variables `id`, `name`, `attack`, `speed`, `defense`, `hp`, `sp_attack`, and `sp_def`.

- Display the first 10 rows of `pok` using `head()` or `slice()`.

Solutions

- Loading `data_pokemon.csv`

```
pok <- read_csv("data_pokemon.csv", show_col_types = FALSE) |>  
  select(id, name, attack, speed, defense, hp, sp_attack, sp_def)
```

- `head()` on `pok`

```
head(pok, n = 10)
```

```
# A tibble: 10 x 8  
  id   name      attack  speed  defense    hp  sp_attack  sp_def  
  <dbl> <chr>     <dbl>  <dbl>   <dbl>  <dbl>     <dbl>  <dbl>  
1 1    Bulbasaur 49     45     49     45     65     65  
2 2    Ivysaur   62     60     63     60     80     80  
3 3    Venusaur  82     80     83     80     100    100  
4 4    Mega Venusaur 100    80    123    80     122    120  
5 5    Charmander 52     65     43     39     60     50  
6 6    Charmeleon 64     80     58     58     80     65  
7 7    Charizard  84     100    78     78     109    85  
8 8    Mega Charizard X 130    100   111    78     130    85  
9 9    Mega Charizard Y 104    100   78     78     159    115  
10 10   Squirtle  48     43     65     44     50     64
```

Question 2. Data management, label variable

Attach descriptive labels to each variable in the dataset pok.
This helps make outputs (e.g., summaries or regression tables) more readable.

Hint: use `relabel()` from the `{collapse}` package.

Solutions

- Add labels to each variable : we use `relabel()` to attach human-friendly names to our variables.

```
pok <- pok |>
  relabel(
    attack = "Attack power", speed = "Speed power", defense = "Defense power",
    hp = "Hit points (health)", sp_attack = "Special attack power",
    sp_def = "Special defense power", id = "ID", name = "Pokemon name"
  )
```

- Check that labels were added: the function `namlab()` (from `{collapse}`) shows variable names, labels, and basic info.

```
namlab(pok, N = TRUE, Ndistinct = TRUE, class = TRUE)
```

| | Variable | Class | N | Ndist | Label |
|---|-----------|-----------|-----|-------|-----------------------|
| 1 | id | numeric | 800 | 800 | ID |
| 2 | name | character | 800 | 800 | Pokemon name |
| 3 | attack | numeric | 800 | 111 | Attack power |
| 4 | speed | numeric | 800 | 108 | Speed power |
| 5 | defense | numeric | 800 | 103 | Defense power |
| 6 | hp | numeric | 800 | 94 | Hit points (health) |
| 7 | sp_attack | numeric | 800 | 105 | Special attack power |
| 8 | sp_def | numeric | 800 | 92 | Special defense power |

- To check only for some variables, we can use `vlabels()` from `{collapse}`

```
vlabels(pok$attack)
```

```
[1] "Attack power"
```

```
vlabels(pok[c("speed", "defense")])
```

```
speed           defense
"Speed power"  "Defense power"
```

Question 3. Summary statistics

For the variables attack, speed, defense, hp, sp_attack, sp_def, compute summary statistics: mean, standard deviation, median, Q1, Q3, minimum, maximum.

Hint: `descr()`, `describe_distribution()`, `get_summary_stats()` `summarise()`, `tbl_summary()`

Solutions

- Save numeric variable names in `numeric_vars`

```
numeric_vars <- names(pok)[-c(1, 2)]  
numeric_vars
```

```
[1] "attack"      "speed"       "defense"     "hp"          "sp_attack"  "sp_def"
```

- Save predictors variable names in `predictors`

```
predictors <- numeric_vars[-1]  
predictors
```

```
[1] "speed"       "defense"     "hp"          "sp_attack"  "sp_def"
```

- Generate summary statistics with `tbl_summary()` from `{gtsummary}`

```
tab_summary <- select(pok, all_of(numeric_vars)) |>
  tbl_summary(
    type = all_continuous() ~ "continuous2",
    statistic = all_continuous() ~ c(
      "{mean} ({sd})", "{median} ({p25}, {p75})", "{min}, {max}"
    ),
    digits = ~ 1
  ) |>
  bold_labels()
```

- Display the table in the PDF using `{kableExtra}`

```
as_kable_extra(tab_summary, booktabs = TRUE, longtable = TRUE, linesep = "") |>
  kable_styling(
    position = "center", font_size = 10,
    latex_options = c("basic", "repeat_header")
  )
```

| Characteristic | N = 800 |
|-----------------------|--------------------|
| Attack power | |
| Mean (SD) | 79.0 (32.5) |
| Median (Q1, Q3) | 75.0 (55.0, 100.0) |
| Min, Max | 5.0, 190.0 |
| Speed power | |
| Mean (SD) | 68.3 (29.1) |
| Median (Q1, Q3) | 65.0 (45.0, 90.0) |
| Min, Max | 5.0, 180.0 |
| Defense power | |
| Mean (SD) | 73.8 (31.2) |
| Median (Q1, Q3) | 70.0 (50.0, 90.0) |
| Min, Max | 5.0, 230.0 |
| Hit points (health) | |
| Mean (SD) | 69.3 (25.5) |
| Median (Q1, Q3) | 65.0 (50.0, 80.0) |
| Min, Max | 1.0, 255.0 |
| Special attack power | |
| Mean (SD) | 72.8 (32.7) |
| Median (Q1, Q3) | 65.0 (49.5, 95.0) |
| Min, Max | 10.0, 194.0 |
| Special defense power | |
| Mean (SD) | 71.9 (27.8) |
| Median (Q1, Q3) | 70.0 (50.0, 90.0) |
| Min, Max | 20.0, 230.0 |

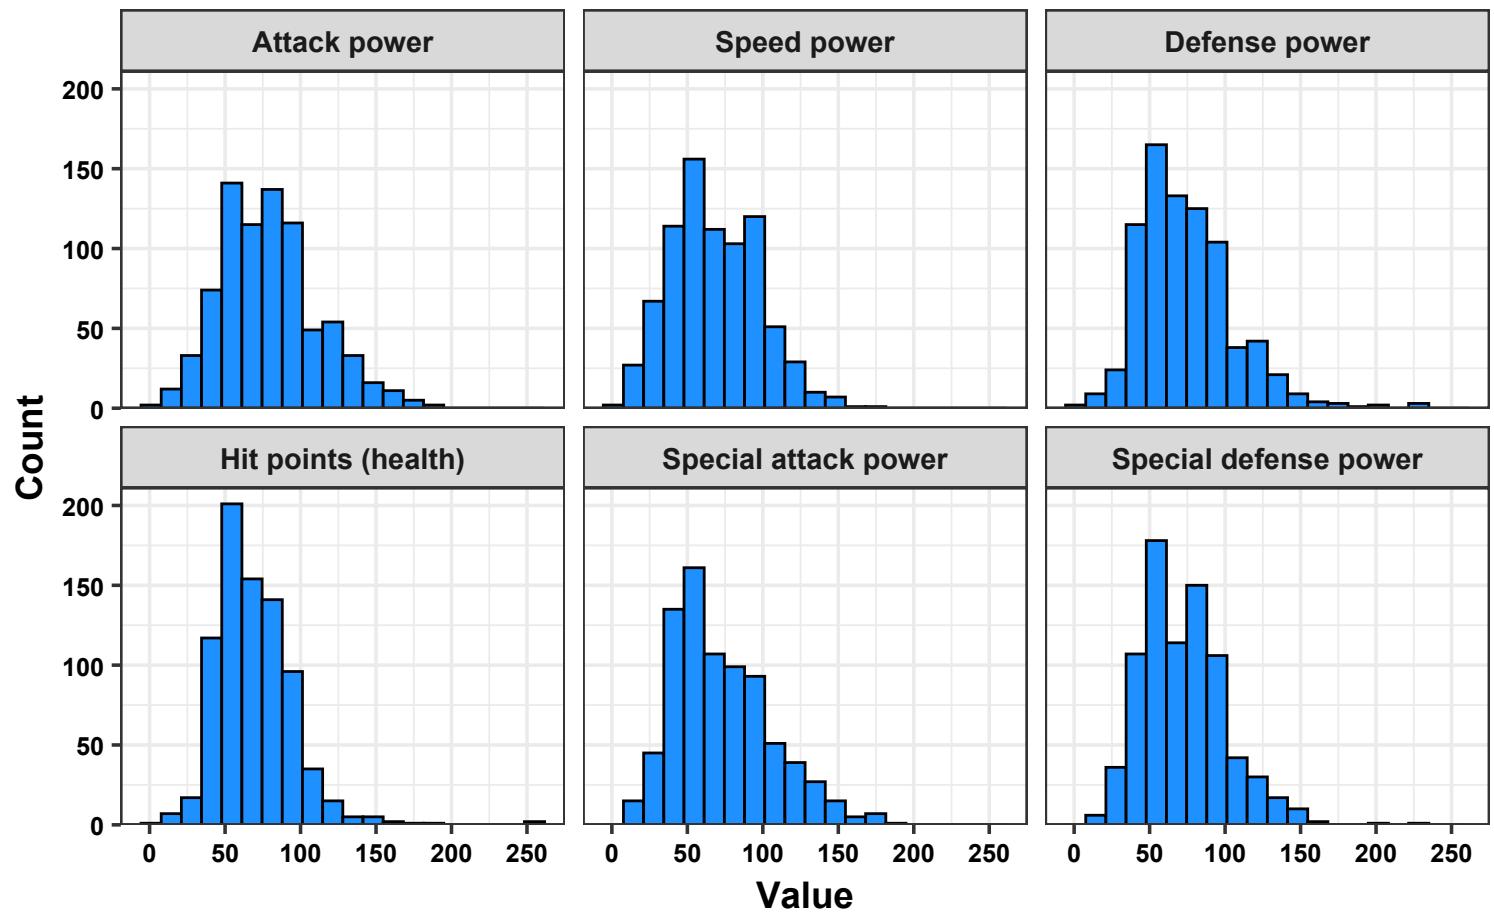
Question 4. Histogram, Scatter plots, Correlations

1. Plot histogram of attack, speed, defense, hp, sp_attack, sp_def.
2. Create scatter plots of attack, speed, defense, hp, sp_attack, sp_def against each other using `ggpairs()` from `{GGally}`.

Solutions

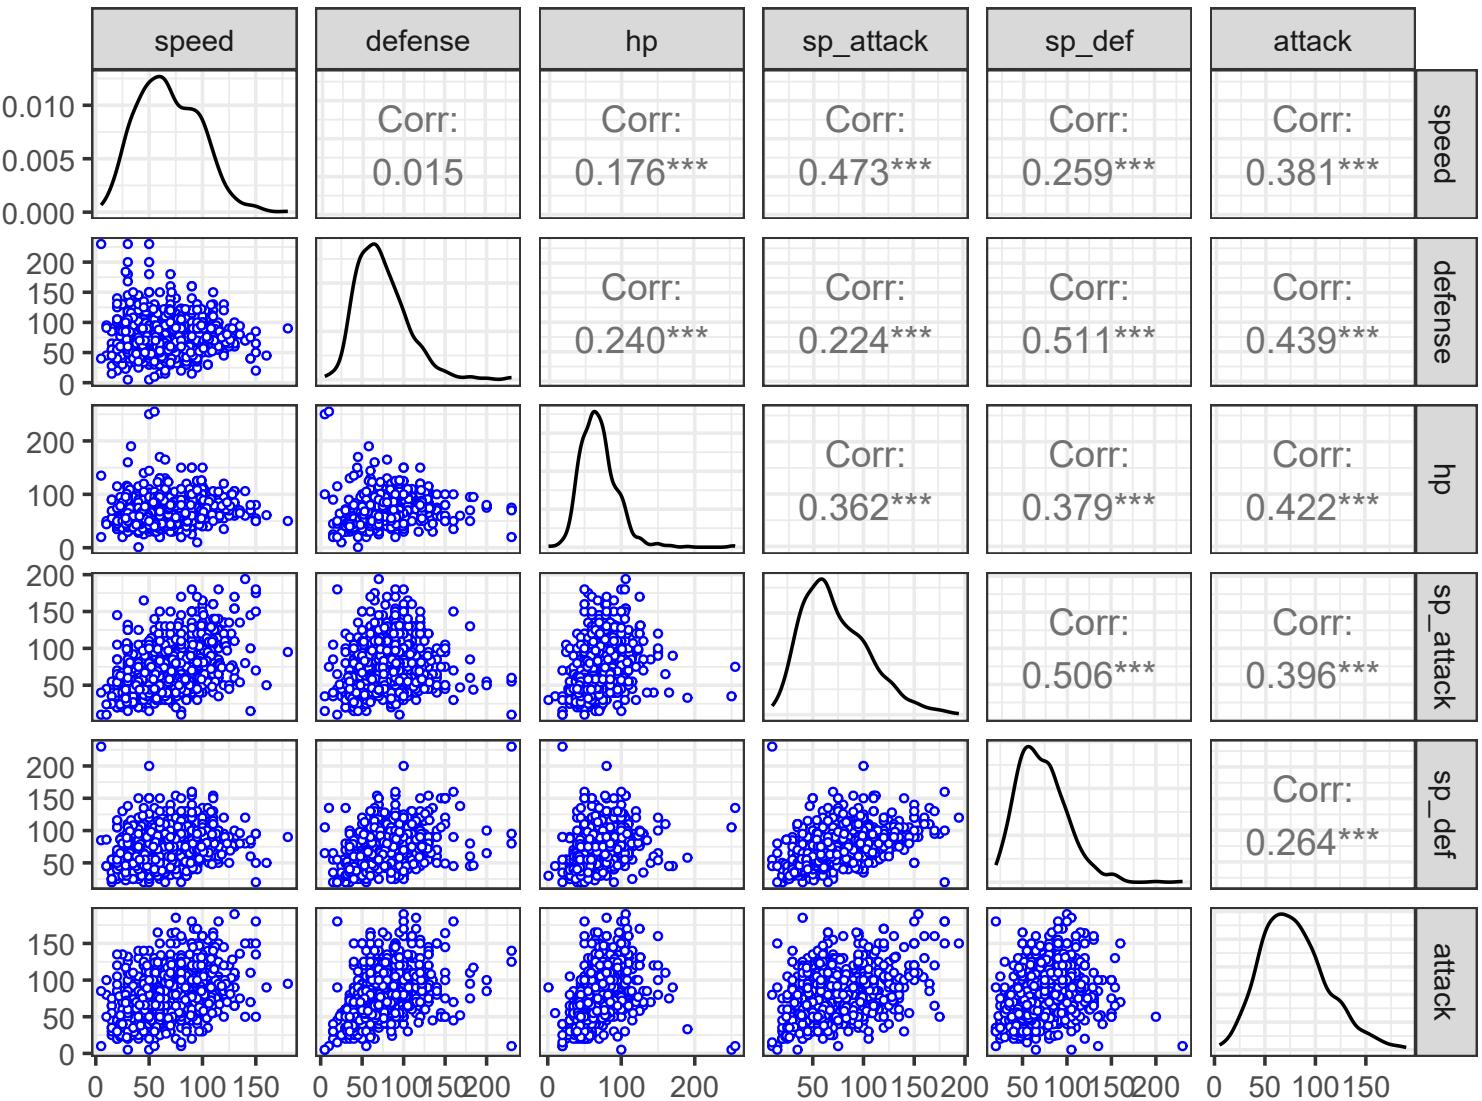
- Histogram of attack, speed, defense, hp, sp_attack, sp_def

```
select(pok, id, attack:sp_def) |> pivot_longer(-id, names_to = "var") |>
  mutate(var = factor(var, levels = numeric_vars, labels = vlabels(pok[, numeric_vars]))) |>
  ggplot(aes(x = value)) +
  facet_wrap(vars(var)) +
  geom_histogram(fill = "dodgerblue", color = "black", bins = 20, linewidth = 0.5) +
  scale_y_continuous(expand = expansion(c(0, 0.05))) +
  scale_x_continuous(breaks = pretty_breaks()) +
  labs(x = "Value", y = "Count") + theme_bw(base_size = 14) +
  theme(strip.text = element_text(size = 11, face = "bold")) +
  labs_pubr()
```



- Scatter plots of attack, speed, defense, hp, sp_attack, sp_def with `ggpairs()`

```
select(pok, -id, -name) |>
  relocate(attack, .after = last_col()) |>
  ggpairs(
    lower = list(
      continuous = wrap(
        "points",
        size = 1, shape = 21, fill = "white", color = "blue", alpha = 1
      )
    )
  ) +
  theme_bw(base_size = 14)
```



Question 5. Multivariate linear gaussian regression model

We now fit a multivariate linear model of `speed`, `defense`, `hp`, `sp_attack`, `sp_def` on `attack`.

Scalar form

$$y_i = \beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip} + \varepsilon_i$$
$$y_i = \mathbf{x}_i^\top \boldsymbol{\beta} + \varepsilon_i \quad p = 5, i = 1, \dots, n, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

with $y_i = \text{attack}_i$ and $(x_{i1}, \dots, x_{i5}) = (\text{speed}_i, \text{defense}_i, \text{hp}_i, \text{sp_attack}_i, \text{sp_def}_i)$

$$\text{attack} = \beta_0 + \beta_1 \text{speed} + \beta_2 \text{defense} + \beta_3 \text{hp} + \beta_4 \text{sp_attack} + \beta_5 \text{sp_def} + \varepsilon$$

Matrix form

$$\mathbf{y} = \mathbb{X}\boldsymbol{\beta} + \varepsilon, \quad \varepsilon \sim \mathcal{N}(\mathbf{0}_n, \sigma^2 \mathbf{I}_n), \text{rank}(\mathbb{X}) = p + 1 = 6$$

1. Fit the model with `lm()` and save the result in `full_model`.
2. Interpret the output of:

```
summary(full_model)
model_parameters(full_model, pretty_names = FALSE)
```

Solutions

- Generate the model formula with `reformulate()`

```
full_formula <- reformulate(term.labels = predictors, response = "attack")
full_formula
```

```
attack ~ speed + defense + hp + sp_attack + sp_def
```

```
class(full_formula)
```

```
[1] "formula"
```

- Fit the full model with `lm()` and save it in `full_model`

```
full_model <- lm(full_formula, data = pok)
```

- Examine the model output

```
summary(full_model)
```

```
Call:
lm(formula = full_formula, data = pok)

Residuals:
    Min      1Q  Median      3Q     Max 
-86.93 -15.90   -2.48  13.50  95.15 

Coefficients:
            Estimate Std. Error t value    Pr(>|t|)    
(Intercept) 3.5187    3.3987   1.04       0.3    
speed        0.3422    0.0340   10.07 < 0.000000000000002 *** 
defense      0.4677    0.0326   14.36 < 0.000000000000002 *** 
hp           0.3700    0.0374   9.88 < 0.000000000000002 *** 
sp_attack    0.1654    0.0342   4.84    0.000001554878 *** 
sp_def       -0.2794   0.0416  -6.71    0.000000000037 *** 
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 

Residual standard error: 24.4 on 794 degrees of freedom
Multiple R-squared:  0.44, Adjusted R-squared:  0.436 
F-statistic: 125 on 5 and 794 DF,  p-value: <0.000000000000002
```

$$\hat{\beta} = (\mathbb{X}^\top \mathbb{X})^{-1} \mathbb{X}^\top \mathbf{y} = (3.519, 0.342, 0.468, 0.370, 0.165, -0.279)^\top$$

```
model_parameters(full_model, pretty_names = FALSE)
```

| Parameter | Coefficient | SE | 95% CI | t(794) | p |
|-------------|-------------|-------|--------------------|--------|-------|
| (Intercept) | 3.519 | 3.399 | [-3.153, 10.190] | 1.035 | 0.301 |
| speed | 0.342 | 0.034 | [0.275, 0.409] | 10.067 | <.001 |
| defense | 0.468 | 0.033 | [0.404, 0.532] | 14.361 | <.001 |
| hp | 0.370 | 0.037 | [0.297, 0.444] | 9.883 | <.001 |
| sp_attack | 0.165 | 0.034 | [0.098, 0.232] | 4.841 | <.001 |
| sp_def | -0.279 | 0.042 | [-0.361, -0.198] | -6.711 | <.001 |

$$\widehat{\text{attack}} = 3.519 + 0.342 \cdot \text{speed} + 0.468 \cdot \text{defense} + 0.37 \cdot \text{hp} + 0.165 \cdot \text{sp_attack} - 0.279 \cdot \text{sp_def}$$

$$\hat{\mathbf{y}} = \mathbb{X} \hat{\beta} = \mathbb{X} (\mathbb{X}^\top \mathbb{X})^{-1} \mathbb{X}^\top \mathbf{y} = P_{\mathbb{X}} \mathbf{y} = H_{\mathbb{X}} \mathbf{y}$$

Interpretation of the regression output

$$\text{attack}_i = \beta_0 + \beta_1 \text{speed}_i + \beta_2 \text{defense}_i + \beta_3 \text{hp}_i + \beta_4 \text{sp_attack}_i + \beta_5 \text{sp_def}_i + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma^2).$$

1. Model fit (global statistics)

$$\hat{\sigma}^2 = \frac{1}{n-r} \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \frac{\hat{\varepsilon}^\top \hat{\varepsilon}}{n-r} = \frac{\|\hat{\varepsilon}\|^2}{n-r} = \frac{\text{RSS}}{n-r}$$

- **Residual standard error (estimated standard deviation of the errors):** $\hat{\sigma} = 24.4$
On average, predictions of attack deviate from the true values by about 24 points.

$$R^2 = \frac{\text{MSS}}{\text{TSS}} = \frac{\sum_{i=1}^n (\hat{y}_i - \bar{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} = 1 - \frac{\text{RSS}}{\text{TSS}} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

- $R^2 = 0.44$ (Adjusted $R^2 = 0.436$)
The model explains about 44% of the variation in attack strength across Pokémons.
- F-statistic = 125 on (5, 794) df, ($p < 2 \times 10^{-16}$)
The model as a whole is highly significant; at least one predictor is associated with attack.

2. Coefficients

- We interpret coefficients while holding other predictors constant.
- **Intercept** ($\hat{\beta}_0 = 3.52, p = 0.30$) Not statistically significant. Represents expected attack when all predictors are 0 (not meaningful here, but needed for the model).
- **Speed** ($\hat{\beta}_1 = 0.34, p < 0.001$) A one-unit increase in speed is associated with a 0.34 increase in attack.
95% CI: [0.28, 0.41].
- **Defense** ($\hat{\beta}_2 = 0.47, p < 0.001$) A one-unit increase in defense is associated with a 0.47 increase in attack. Strongest positive effect, CI : [0.40, 0.53].
- **HP** ($\hat{\beta}_3 = 0.37, p < 0.001$) A one-unit increase in HP is associated with a 0.37 increase in attack.
CI : [0.30, 0.44].
- **Special Attack** ($\hat{\beta}_4 = 0.17, p < 0.001$) A one-unit increase in special attack is associated with a 0.17 increase in attack. CI : [0.10, 0.23].
- **Special Defense** ($\hat{\beta}_5 = -0.28, p < 0.001$) A one-unit increase in special defense is associated with a 0.28 decrease in attack. CI : [-0.36, -0.20].
Suggests a tradeoff: high special defense Pokémons tend to have weaker physical attack.

3. Substantive interpretation

- Pokémon with higher defense, HP, and speed tend to also have higher attack.
- Special attack contributes positively, but less strongly.
- Special defense shows an *inverse relationship*: more defensive Pokémon are less offensively strong.

Conclusion:

The regression explains about 44% of the variance in attack strength. Most predictors are significant and meaningful: defense, HP, speed, and special attack increase attack power, while special defense decreases it.

Question 6. Test of overall regression

We want to test

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_5 = 0 \quad \text{versus} \quad H_1 : \text{At least one } \beta_j \neq 0$$
$$\Leftrightarrow H_0 : (m_0) \mathbf{y} = \beta_0 \mathbf{1}_n + \boldsymbol{\varepsilon} \quad \text{versus} \quad H_1 : (m_1) \mathbf{y} = \mathbb{X}\boldsymbol{\beta} + \boldsymbol{\varepsilon}$$

Perform this test in 2 different ways. Hint: Fisher test for nested models / General linear hypothesis tests

Method 1: Fisher test for nested models

The F-statistic is

$$F = \frac{\left\| P_{m_0} \mathbf{y} - P_{m_1} \mathbf{y} \right\|^2 / (p - q)}{\left\| \mathbf{y} - P_{m_1} \mathbf{y} \right\|^2 / (n - r)} = \frac{[\text{RSS}(m_0) - \text{RSS}(m_1)] / (p - q)}{\text{RSS}(m_1) / (n - r)} \sim F_{p-q, n-r} \quad (\text{under } H_0)$$

- Here $q = 0$ (reduced model has only an intercept),
 - $r = p + 1 = 6$,
 - $p - q = 5$.
-

Method 2: General linear hypothesis test

We can also write

$$H_0 : \beta_1 = \beta_2 = \dots = \beta_5 = 0 \quad \text{versus} \quad H_1 : \text{At least one } \beta_j \neq 0$$
$$\Leftrightarrow H_0 : \mathbf{C}\boldsymbol{\beta} = \mathbf{0}_5 \quad \text{versus} \quad H_1 : \mathbf{C}\boldsymbol{\beta} \neq \mathbf{0}_5$$

It is clear that

$$\mathbf{C}\boldsymbol{\beta} = \mathbf{0}_5 \Leftrightarrow \begin{bmatrix} 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \\ \beta_2 \\ \beta_3 \\ \beta_4 \\ \beta_5 \end{bmatrix} = \begin{bmatrix} \beta_1 \\ \beta_2 \\ \beta_3 \\ \beta_4 \\ \beta_5 \end{bmatrix} = \mathbf{0}_5$$

The corresponding test statistic is

$$F = \frac{(\mathbf{C}\hat{\boldsymbol{\beta}})^\top [\mathbf{C}(\mathbb{X}^\top \mathbb{X})^{-1} \mathbf{C}^\top]^{-1} (\mathbf{C}\hat{\boldsymbol{\beta}}) / q}{\hat{\sigma}^2} \sim F_{q, n-r}, \quad q = 5.$$

Solutions

Method 1: Nested model comparison

- Fit the reduced model (intercept only)

```
null_model <- lm(attack ~ 1, data = pok)
```

- Perform the overall test with `anova()`

```
global_test <- anova(null_model, full_model)
global_test
```

Analysis of Variance Table

```
Model 1: attack ~ 1
Model 2: attack ~ speed + defense + hp + sp_attack + sp_def
  Res.Df   RSS Df Sum of Sq    F    Pr(>F)
1     799 841731
2     794 471595  5     370136 125 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
global_test$F
```

```
[1] NA 124.64
```

```
summary(full_model)[["fstatistic"]]
```

```
value numdf dendf
124.64  5.00 794.00
```

Method 2: General linear hypothesis test

- Compute F statistic with `linearHypothesis()` from `{car}`

$$F = \frac{(\mathbf{C}\hat{\beta})^\top [\mathbf{C}(\mathbb{X}^\top \mathbb{X})^{-1} \mathbf{C}^\top]^{-1} (\mathbf{C}\hat{\beta}) / q}{\hat{\sigma}^2}$$

`predictors`

```
[1] "speed"      "defense"     "hp"          "sp_attack"   "sp_def"
```

```
linearHypothesis(full_model, predictors)
```

```
Linear hypothesis test:
speed = 0
defense = 0
hp = 0
sp_attack = 0
sp_def = 0

Model 1: restricted model
Model 2: attack ~ speed + defense + hp + sp_attack + sp_def

Res.Df    RSS Df Sum of Sq    F            Pr(>F)
1    799 841731
2    794 471595  5    370136 125 <0.0000000000000002 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- We can also compute F statistic manually

$$F = \frac{(\mathbf{C}\hat{\beta})^\top [\mathbf{C}(\mathbb{X}^\top \mathbb{X})^{-1} \mathbf{C}^\top]^{-1} (\mathbf{C}\hat{\beta}) / q}{\hat{\sigma}^2}$$

```
C <- cbind(rep(0, 5), diag(nrow = 5))
betas <- coef(full_model)

# Design matrix
X <- model.matrix(full_model)

# Residual variance estimate ( $\sigma^2$ -hat)
sigma2_hat <- sigma(full_model)^2

# Number of restrictions
q <- nrow(C)

# Numerator: ( $\mathbf{C} \beta_{\text{hat}}$ )
C_beta <- C %*% betas

# Middle matrix: [ $\mathbf{C} (\mathbf{X}'\mathbf{X})^{-1} \mathbf{C}'$ ]^{-1}
middle <- solve(C %*% solve(t(X) %*% X) %*% t(C))

# F-statistic
F_stat <- as.numeric(t(C_beta) %*% middle %*% C_beta / (q * sigma2_hat))
```

- F-statistic

```
F_stat
```

```
[1] 124.64
```

- Critical value and p-value

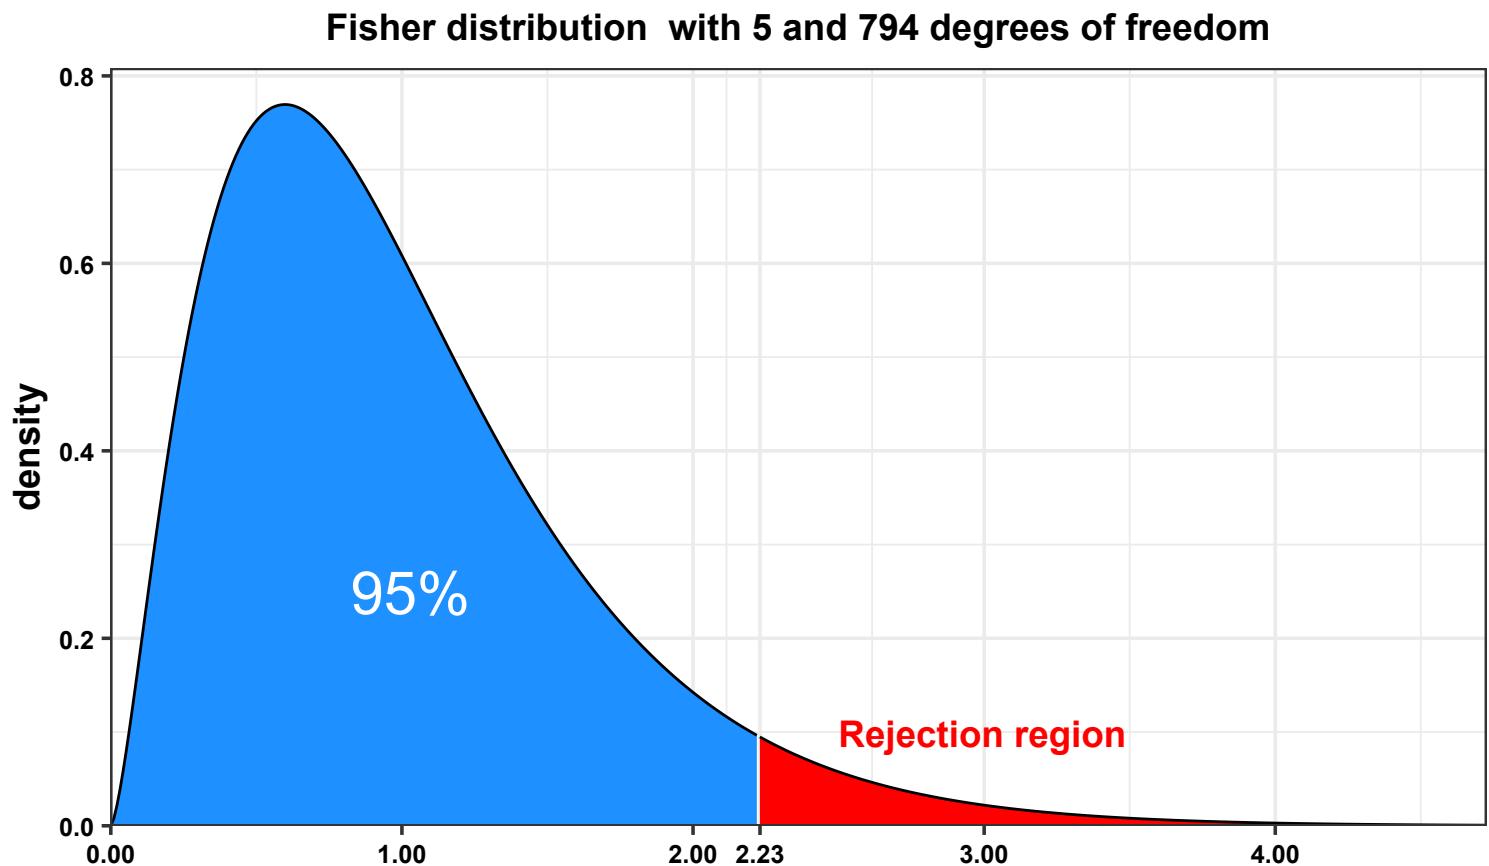
```
# 95% quantile of  $F(q, n-r)$ 
qf(0.95, q, df.residual(full_model), lower.tail = TRUE)
```

```
[1] 2.2254
```

```
# p-value
pf(F_stat, q, df.residual(full_model), lower.tail = FALSE) |> label_scientific()()
```

```
[1] "2.27e-97"
```

- Visualize rejection region of the Fisher distribution



Question 7. Indices of model performance for regression

Compute indices of performance for the `full_model`. Hint: `glance()`, `model_performance()`

Solutions

- With `glance()` from `{broom}`

```
glance(full_model)
```

```
# A tibble: 1 x 12
  r.squared adj.r.squared sigma statistic p.value    df logLik   AIC   BIC deviance df.residual
    <dbl>        <dbl> <dbl>      <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>        <int>
1     0.440       0.436  24.4     125.  2.27e-97      5 -3687. 7388. 7421.  471595.      794
```

$$R^2 = \frac{\text{MSS}}{\text{TSS}} = 1 - \frac{\text{RSS}}{\text{TSS}} \quad R_a^2 = 1 - \frac{(n-1)\text{RSS}}{(n-m)\text{TSS}} \quad (m = p+1)$$

$$\text{AIC}(m) = n \log\left(\frac{\text{RSS}(m)}{n}\right) + 2m \quad \text{BIC}(m) = n \log\left(\frac{\text{RSS}(m)}{n}\right) + \log(n) \times m$$

$$\text{deviance}(m) = \text{RSS}(m) \quad \text{df.residual}(m) = n - m \quad \text{sigma}(m) = \sqrt{\frac{\text{RSS}}{n-m}} = \hat{\sigma}$$

- With `model_performance()` from `{performance}`

```
model_performance(full_model)
```

```
# Indices of model performance
  AIC |  AICc |  BIC |  R2 | R2 (adj.) |  RMSE |  Sigma
-----+
7387.7 | 7387.9 | 7420.5 | 0.440 |      0.436 | 24.279 | 24.371
```

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

- Other useful functions

```
AIC(full_model)
```

```
[1] 7387.7
```

```
BIC(full_model)
```

```
[1] 7420.5
```

```
r2(full_model, ci = 0.95) # {performance}
```

```
R2: 0.440 [0.384, 0.487]  
adj. R2: 0.436 [0.380, 0.484]
```

```
summary(full_model)[["r.squared"]]
```

```
[1] 0.43973
```

```
summary(full_model)[["adj.r.squared"]]
```

```
[1] 0.4362
```

Question 8. Joint hypothesis test

Consider the full regression model `full_model`

$$\text{attack} = \beta_0 + \beta_1 \text{speed} + \beta_2 \text{defense} + \beta_3 \text{hp} + \beta_4 \text{sp_attack} + \beta_5 \text{sp_def} + \varepsilon$$

We want to test jointly whether the coefficients on `sp_attack` and `sp_def` are equal to zero:

$$H_0 : \beta_4 = \beta_5 = 0 \quad \text{versus} \quad H_1 : \text{at least one of } \beta_4, \beta_5 \text{ is nonzero.}$$

Hints:

- Compare the full model with a restricted model (without `sp_attack` and `sp_def`) using an F-test (`anova()`).
- Use a joint Wald test (`linearHypothesis()`, `waldtest()`).

Solutions

Method 1: Nested model comparison

- Fit the restricted model, `res_model`

```
res_model <- lm(attack ~ speed + defense + hp, data = pok)
```

- Compare with the full model using `anova()`:

```
anova(res_model, full_model) |> qTBL()
```

```
# A tibble: 2 x 6
  Res.Df    RSS    Df `Sum of Sq`      F    `Pr(>F)`
  <dbl>   <dbl>  <dbl>      <dbl> <dbl>    <dbl>
1    796 502738.     NA        NA    NA    NA
2    794 471595.     2     31143.  26.2 9.42e-12
```

- Compute F statistics by hand: $F = \frac{[\text{RSS}(m_0) - \text{RSS}(m_1)] / (p-q)}{\text{RSS}(m_1) / (n-r)} \sim F_{p-q, n-r}$ (under H_0)

```
rss0 <- deviance(res_model)
df0 <- df.residual(res_model)
rss1 <- deviance(full_model)
df1 <- df.residual(full_model)
fstat <- ((rss0 - rss1) / (df0 - df1)) / (rss1 / df1)
fstat
```

```
[1] 26.217
```

- Compute p-value

```
pf(fstat, df0 - df1, df1, lower.tail = FALSE)
```

```
[1] 0.000000000094234
```

Method 2: Wald-type tests

- Use `linearHypothesis()` from `{car}`

```
linearHypothesis(full_model, c("sp_attack = 0", "sp_def = 0"))
```

```
Linear hypothesis test:
sp_attack = 0
sp_def = 0

Model 1: restricted model
Model 2: attack ~ speed + defense + hp + sp_attack + sp_def

Res.Df   RSS Df Sum of Sq    F      Pr(>F)
1     796 502738
2     794 471595  2     31143 26.2 0.000000000094 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

- Or use `waldtest()` from `{lmtest}`

```
# Compare restricted vs full model
waldtest(res_model, full_model, test = "F")

# Test restrictions directly
waldtest(full_model, c("sp_attack", "sp_def"), test = "F")
```

```
Wald test

Model 1: attack ~ speed + defense + hp + sp_attack + sp_def
Model 2: attack ~ speed + defense + hp
Res.Df Df    F      Pr(>F)
1     794
2     796 -2 26.2 0.000000000094 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Question 9. Prediction & Intervals

Consider `full_model`

1. Compute 95% Confidence Interval for the mean $\mathbb{E}(\text{attack})$ given $\text{speed} = 30, 70, 110, 150$ and fixing the other predictors at their mean

Hint: `predict(..., interval = "confidence")`, `estimate_expectation()`

2. Suppose a new Pokémon is created with the following characteristics :

| speed | defense | hp | sp_attack | sp_def |
|-------|---------|-----|-----------|--------|
| 50 | 42 | 100 | 135 | 60 |

Predict the attack for Pokémon and the appropriate 95%CI.

Solutions

- Create grid of values: $\text{speed} = 30, 70, 110, 150$ and fixing the other predictors at their mean with `get_datagrid()` from `{insight}`

```
grid1 <- select(pok, all_of(predictors)) |>  
  get_datagrid(by = "speed = seq(30, 150, 40)", numerics = "integer")
```

```
grid1
```

Visualisation Grid

| speed | defense | hp | sp_attack | sp_def |
|-------|---------|----|-----------|--------|
| 30 | 74 | 69 | 73 | 72 |
| 70 | 74 | 69 | 73 | 72 |
| 110 | 74 | 69 | 73 | 72 |
| 150 | 74 | 69 | 73 | 72 |

Maintained constant: `defense`, `hp`, `sp_attack`, `sp_def`

- Predicted and 95% Confidence Interval for the prediction

```
estimate_expectation(full_model, data = grid1, ci = 0.95)
```

Model-based Predictions

| speed | Predicted | SE | 95% CI |
|-------|-----------|------|------------------|
| 30 | 65.88 | 1.56 | [62.82, 68.95] |
| 70 | 79.57 | 0.86 | [77.88, 81.27] |
| 110 | 93.26 | 1.66 | [90.00, 96.51] |
| 150 | 106.95 | 2.91 | [101.24, 112.65] |

Variable predicted: attack

Predictors modulated: speed = seq(30, 150, 40)

Predictors controlled: defense (74), hp (69), sp_attack (73), sp_def (72)

- We now predict the attack of a Pokémon with the following characteristics

| speed | defense | hp | sp_attack | sp_def |
|-------|---------|-----|-----------|--------|
| 50 | 42 | 100 | 135 | 60 |

```
grid2 <- c(speed = 50, defense = 42, hp = 100, sp_attack = 135, sp_def = 60) |>
  as_tibble_row()
```

grid2

```
# A tibble: 1 x 5
  speed defense    hp sp_attack sp_def
  <dbl>   <dbl> <dbl>     <dbl>   <dbl>
1     50      42    100      135      60
```

- Prediction interval with `estimate_prediction()`

```
estimate_prediction(full_model, data = grid2, ci = 0.95)
```

Model-based Predictions

| speed | defense | hp | sp_attack | sp_def | Predicted | SE | 95% CI |
|-------|---------|-----|-----------|--------|-----------|-------|-----------------|
| 50 | 42 | 100 | 135 | 60 | 82.84 | 24.56 | [34.63, 131.04] |

Variable predicted: attack

Question 10. Residual diagnostics

Note: Standardized vs Studentized residuals

- Let denote by h_{ij} the element of the projector $P_{\mathbb{X}} = H_{\mathbb{X}}$ such that $P_{\mathbb{X}} = H_{\mathbb{X}} = [h_{ij}]$
- The diagonal elements $h_{ii} \in [0, 1]$ are called the *leverages*
- If $h_{ii} > 2p/n$ (sometimes $h_{ii} > 3p/n$), then the observation i is consider an *outlier*
- Standardized residuals (from `rstandard()`)
Raw residuals are rescaled by their estimated standard deviation, taking into account leverage.

$$\hat{r}_i = \frac{\hat{\varepsilon}_i}{\hat{\sigma} \sqrt{1 - h_{ii}}}$$

where $\hat{\varepsilon}_i$ is the raw residual and h_{ii} is the leverage of observation i .
These make residuals roughly comparable across observations.

- Studentized residuals (from `rstudent()`)
Go one step further: each residual is scaled using a variance estimate that excludes the i -th observation.
This gives more accurate standard errors and makes large outliers easier to detect.

$$t_i^* = \frac{\hat{\varepsilon}_i}{\hat{\sigma}_{(-i)} \sqrt{1 - h_{ii}}}$$

where $\hat{\sigma}_{(-i)}$ is the error standard deviation estimated without observation i .

Using `full_model` and functions from the file `helper_functions.R`:

1. Plot residuals vs fitted values and vs each predictor speed, defense, hp, sp_attack, sp_def.
2. Plot $\sqrt{|\text{Standardized residuals}|}$ vs fitted values and vs each predictor.
3. Plot studentized residuals vs fitted values and vs each predictor.
4. Plot residuals in the order of observation (to detect dependence).
5. Plot a histogram of the standardized residuals.
6. Perform a normality test on standardized residuals.
7. Plot a normal Q-Q plot of standardized residuals.
8. Perform the Breusch–Pagan test for heteroskedasticity.
9. Perform the Durbin–Watson test on the residuals.

Solutions

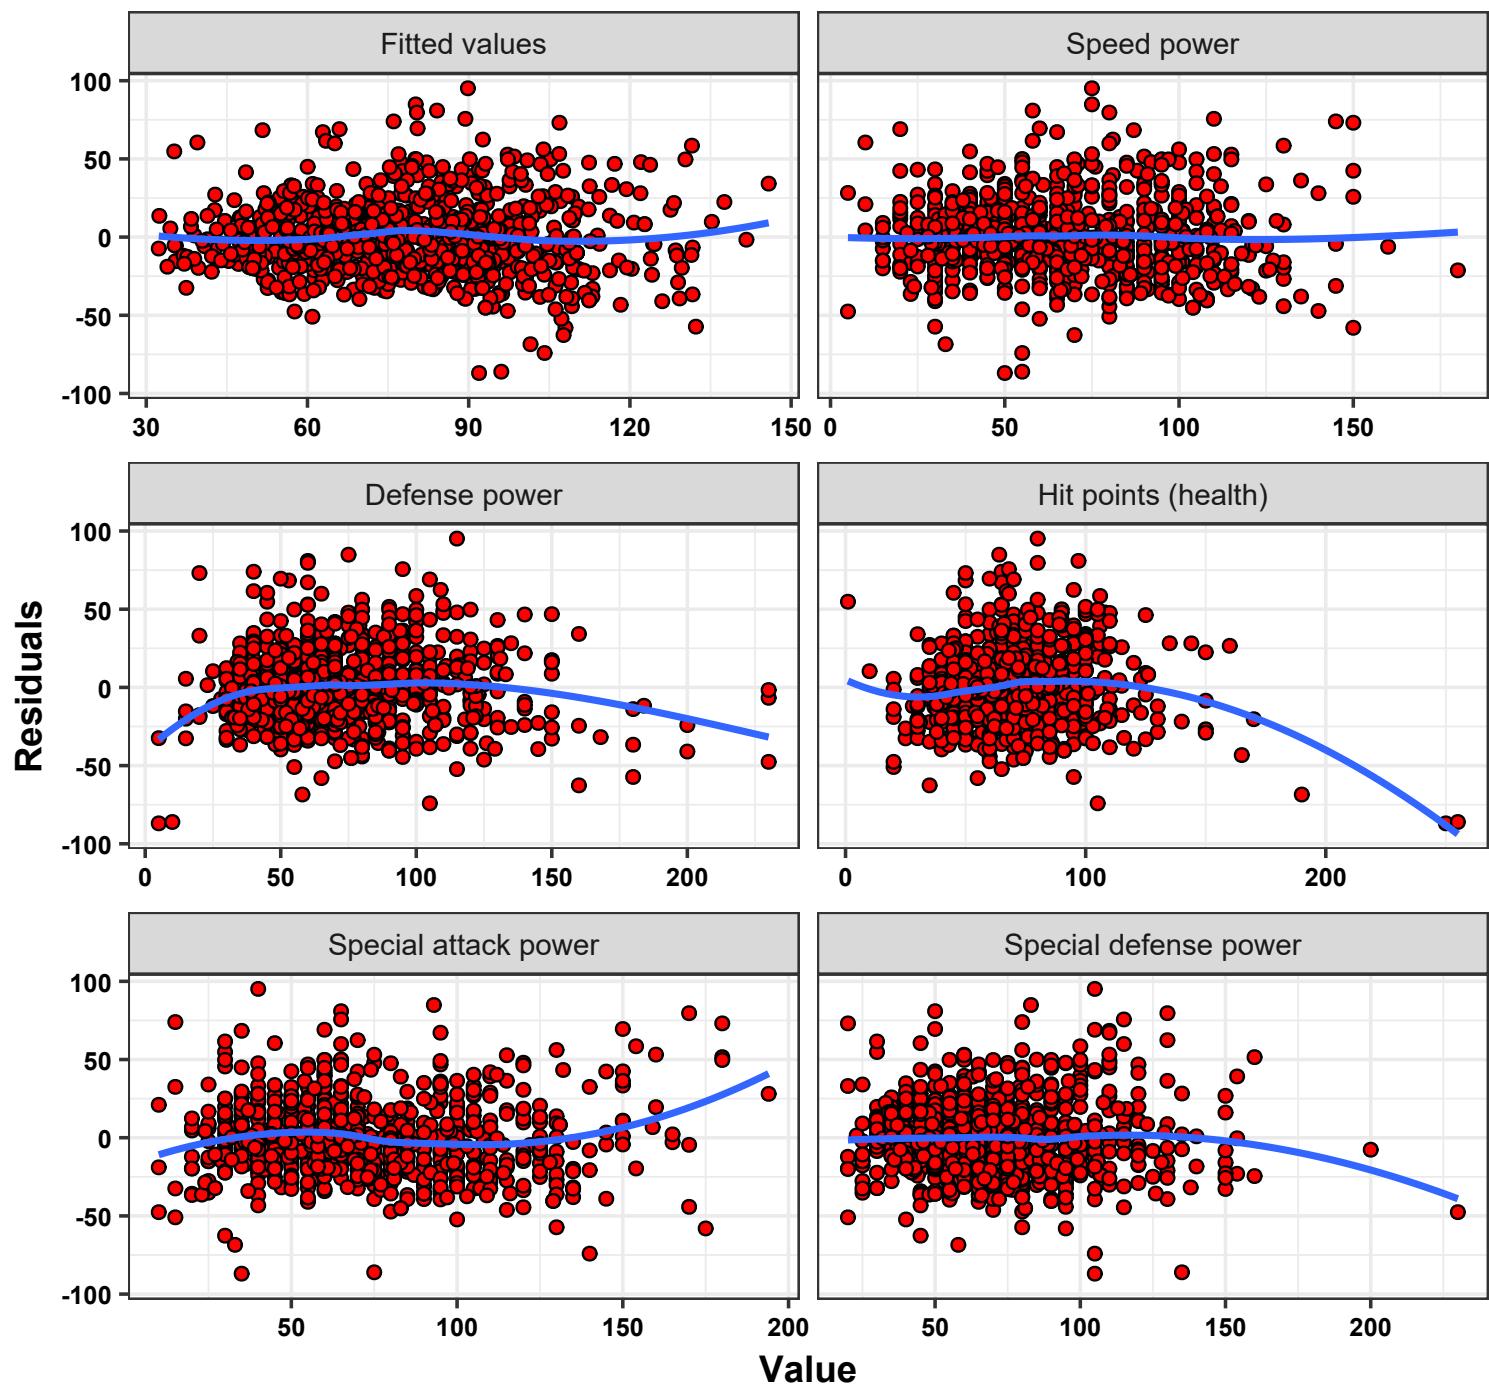
- Prepare residual diagnostics dataset

```
augment_full <- augment(full_model)
head(augment_full)
```

| | attack | speed | defense | hp | sp_attack | sp_def | .fitted | .resid | .hat | .sigma | .cooksdi | .std.resid |
|---|--------|-------|---------|-------|-----------|--------|---------|--------|---------|--------|------------|------------|
| | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> | <dbl> |
| 1 | 49 | 45 | 49 | 45 | 65 | 65 | 51.1 | -2.08 | 0.00404 | 24.4 | 0.00000492 | -0.0854 |
| 2 | 62 | 60 | 63 | 60 | 80 | 80 | 66.6 | -4.60 | 0.00244 | 24.4 | 0.0000145 | -0.189 |
| 3 | 82 | 80 | 83 | 80 | 100 | 100 | 87.9 | -5.92 | 0.00277 | 24.4 | 0.0000274 | -0.243 |
| 4 | 100 | 80 | 123 | 80 | 122 | 120 | 105. | -4.67 | 0.00691 | 24.4 | 0.0000430 | -0.192 |
| 5 | 52 | 65 | 43 | 39 | 60 | 50 | 56.3 | -4.26 | 0.00371 | 24.4 | 0.0000190 | -0.175 |
| 6 | 64 | 80 | 58 | 58 | 80 | 65 | 74.6 | -10.6 | 0.00211 | 24.4 | 0.0000663 | -0.433 |

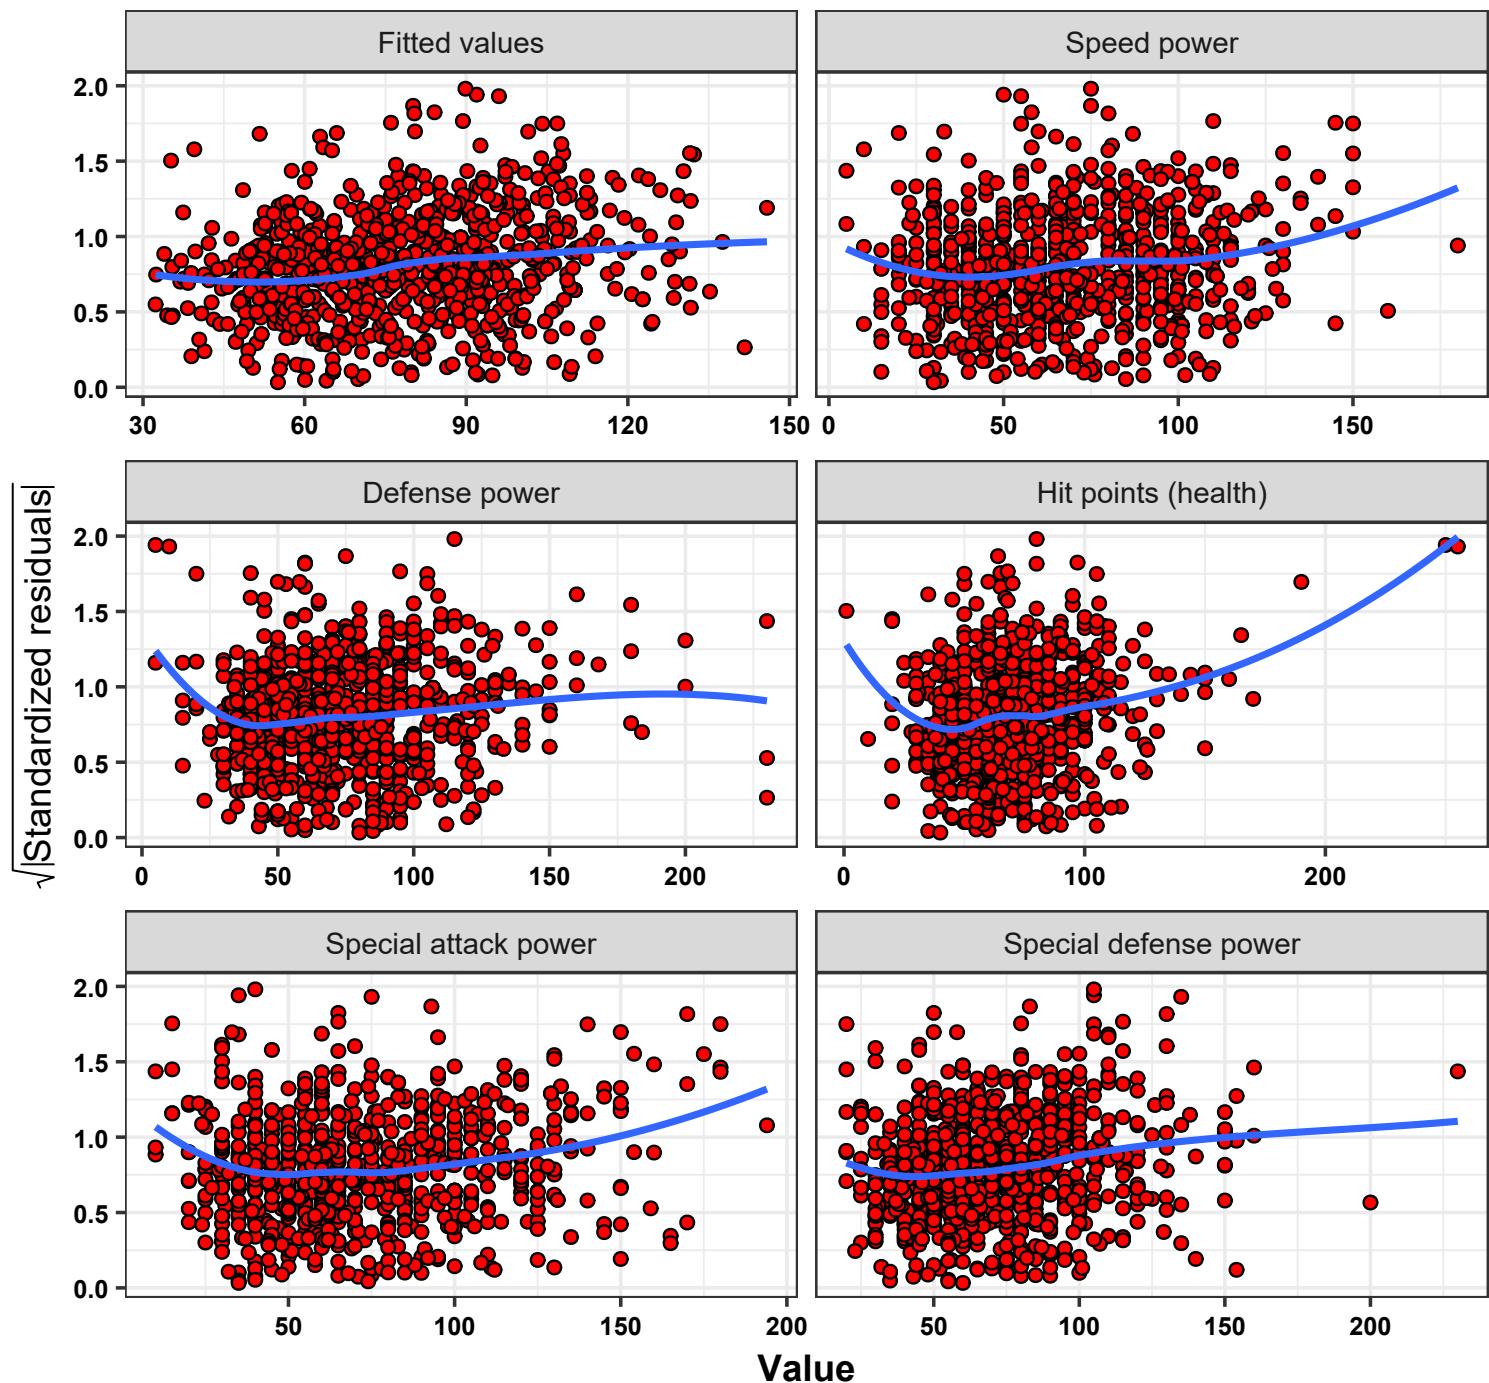
- Residuals vs fitted values and predictors

```
resid_vs_predictors(model = full_model, predictors = predictors)
```



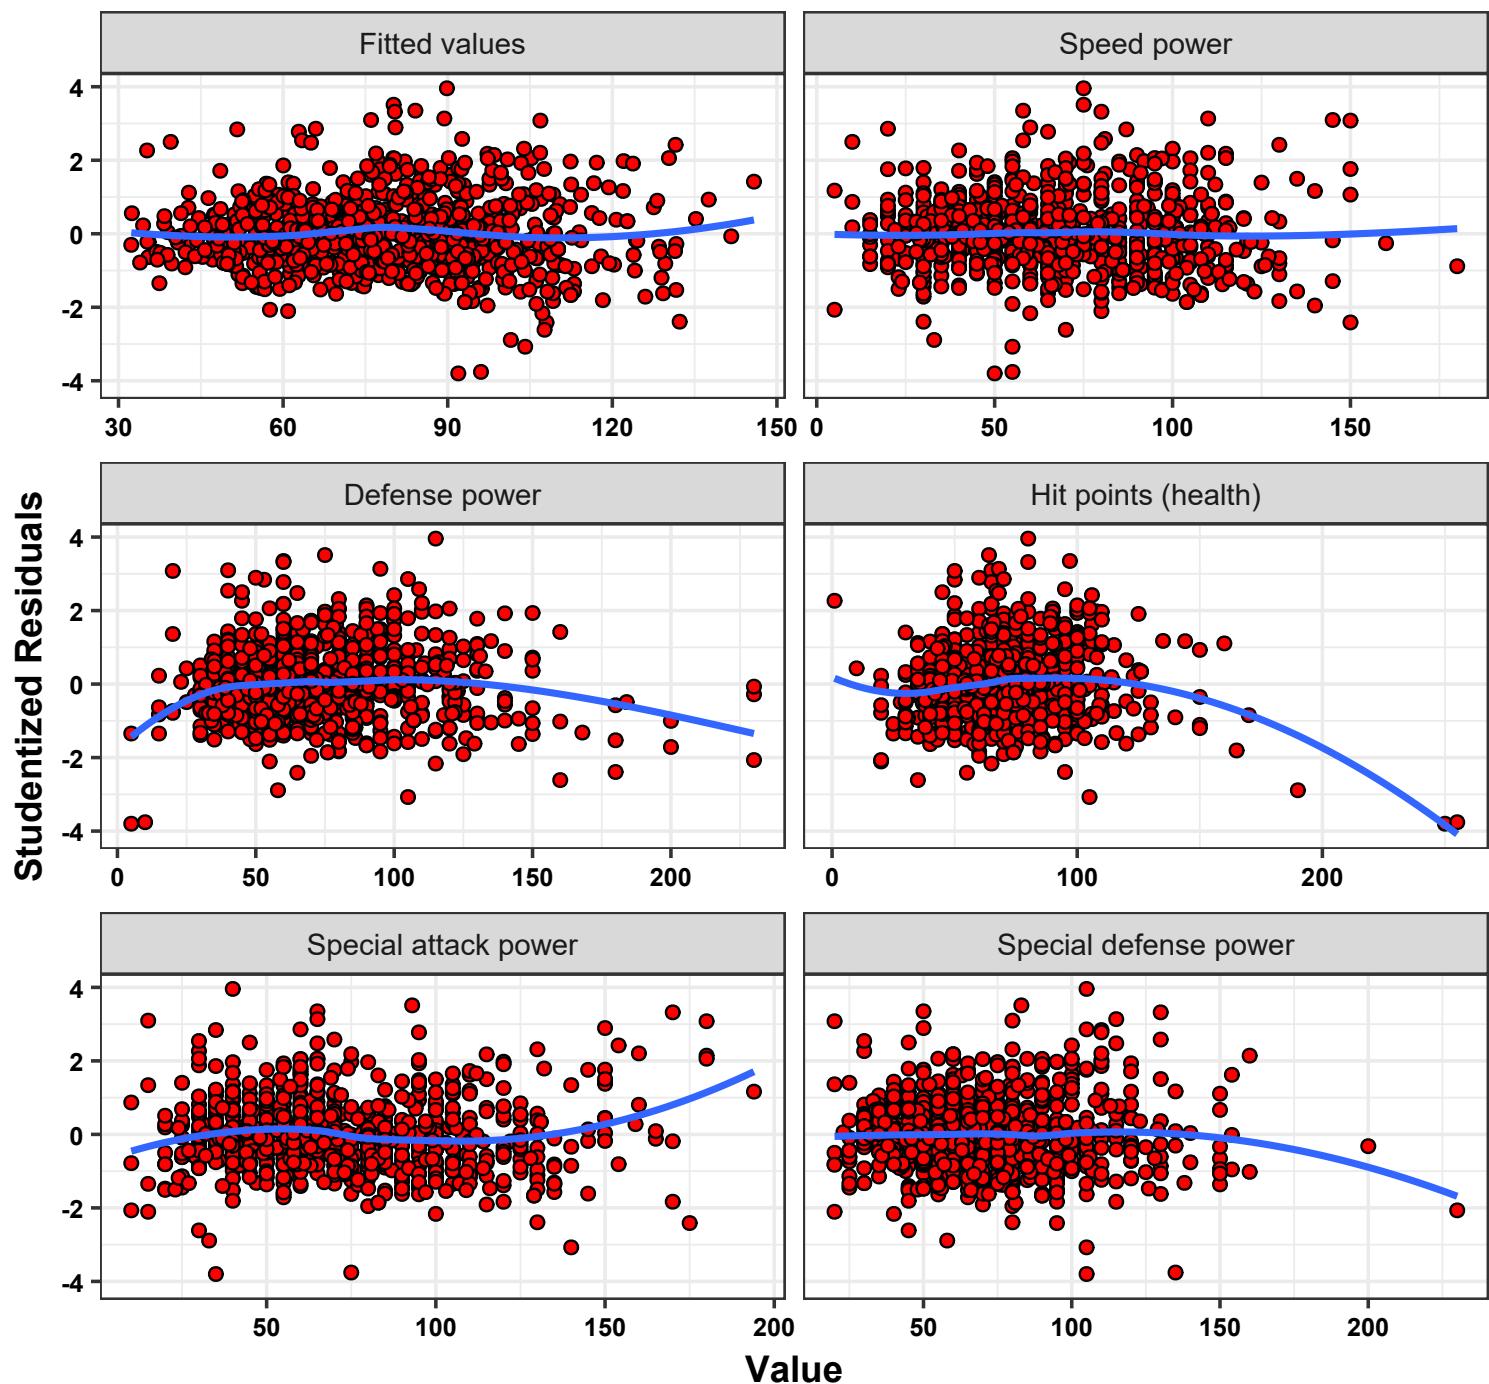
- $\sqrt{|\text{Standardized residuals}|}$ vs fitted values and predictors

```
resid_stand_vs_predictors(model = full_model, predictors = predictors)
```



- Studentized residuals vs fitted values and predictors

```
resid_stud_vs_predictors(full_model, predictors)
```



- Residuals vs observation order

```
p1 <- resid_vs_order(full_model)
```

- Histogram of standardized residuals

```
p2 <- resid_stand_hist(full_model)
```

- Density of standardized residuals

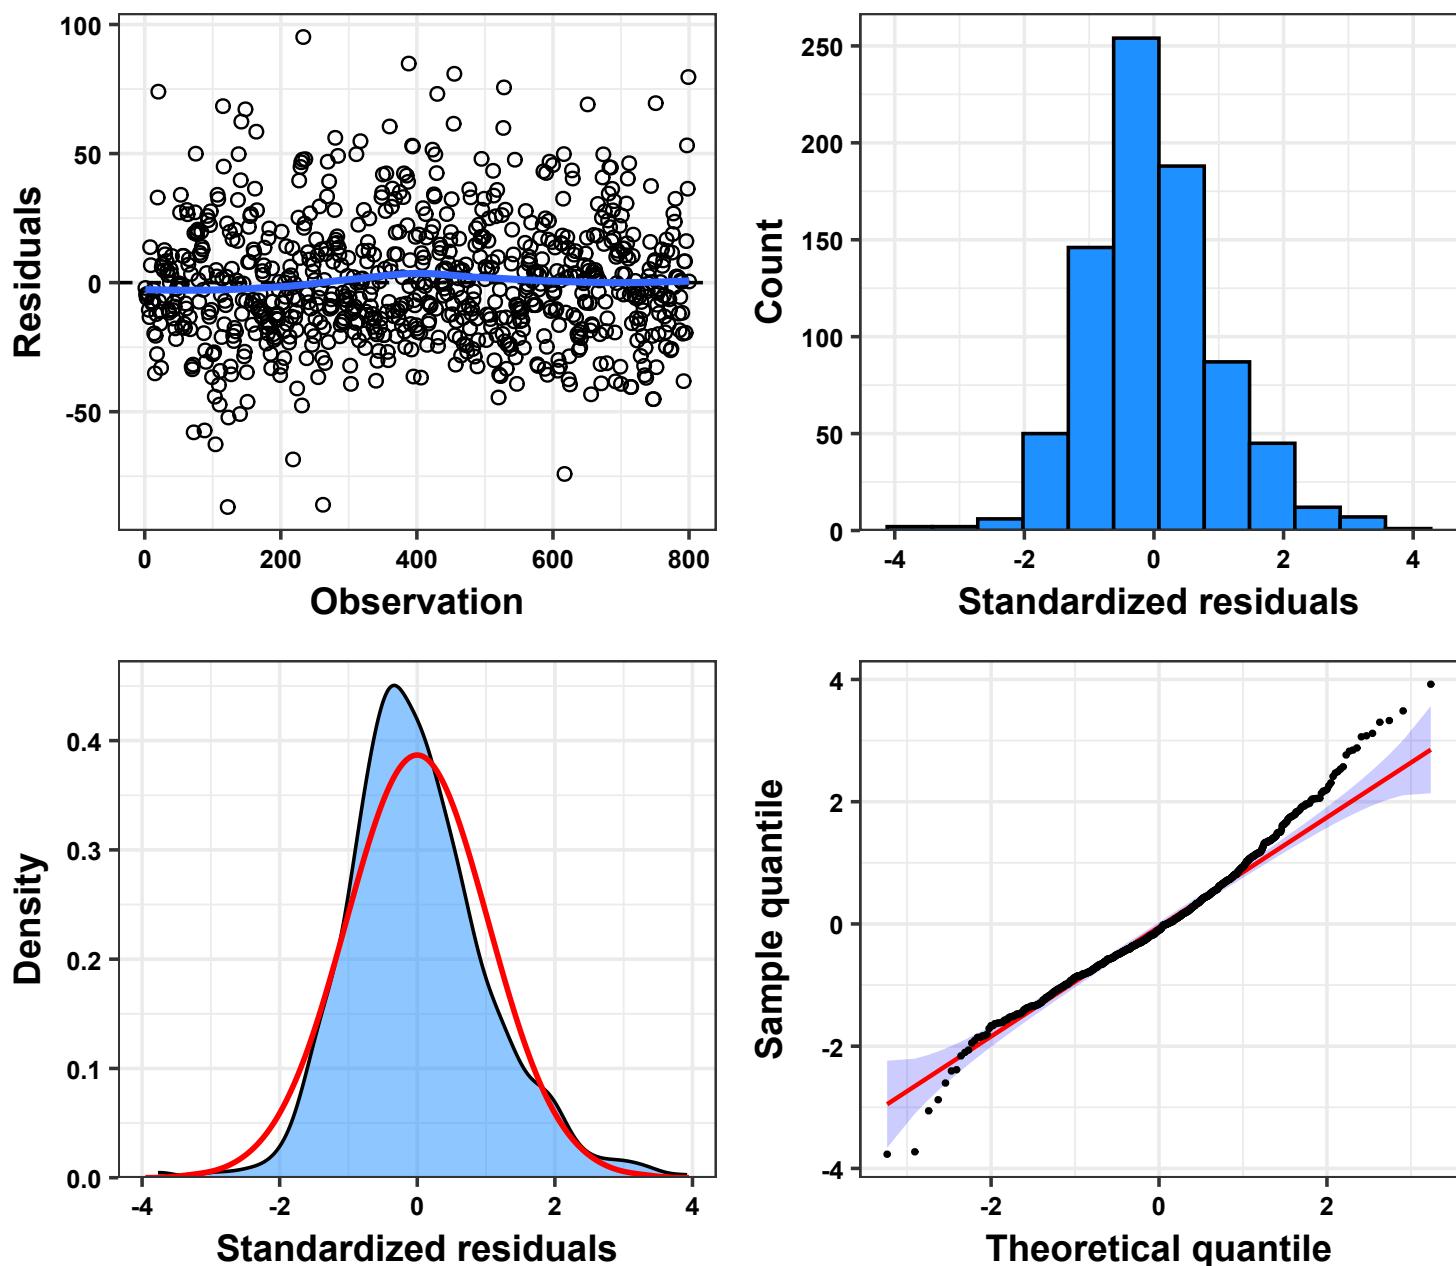
```
p3 <- resid_stand_dens(full_model)
```

- Normal Q-Q plot

```
p4 <- resid_stand_qq(full_model)
```

- We combine the 4 plots with the help of `{patchwork}`

`p1 + p2 + p3 + p4`



- Breusch–Pagan test for heteroskedasticity

```
ncvTest(full_model)
```

```
Non-constant Variance Score Test  
Variance formula: ~ fitted.values  
Chisquare = 37.223, Df = 1, p = 0.00000000105
```

The Breusch–Pagan test strongly rejects the null hypothesis ($p < 10^{-6}$), but with a large sample, even small deviations from perfect homoscedasticity will appear “significant.”

From the residual and scale–location plots, variance seems roughly constant with only mild increases for certain predictors (e.g., `hp`, `sp_def`).

Conclusion:

The test result likely overstates the problem. There is *mild heteroscedasticity*, but not enough to invalidate the model.

- Durbin–Watson test

```
durbinWatsonTest(full_model)
```

```
lag Autocorrelation D-W Statistic p-value  
1          0.30572      1.3886      0  
Alternative hypothesis: rho != 0
```

The Durbin–Watson test gave a statistic of about 1.39 with a very small p -value, suggesting positive autocorrelation.

However, with ($n \approx 800$), the test becomes *too powerful* and flags even trivial correlations as significant.

Moreover, this dataset is not a time series (observations are not ordered chronologically), so the detected autocorrelation may reflect mild structural grouping rather than temporal dependence.

Conclusion:

Although the test is significant, there is no visible pattern in the residual plots.

Residuals appear *approximately independent*, and the practical impact on OLS validity is minimal.

- Shapiro–Wilk test of residual normality

```
augment(full_model) |>
  shapiro_test(.std.resid)
```

```
# A tibble: 1 × 3
  variable    statistic      p
  <chr>        <dbl>      <dbl>
1 .std.resid   0.981 0.0000000102
```

The Shapiro–Wilk test produced a tiny p -value ($p < 10^{-7}$), but with large n , it detects even negligible deviations.

The histogram and density of standardized residuals appear symmetric and bell-shaped, and the Q–Q plot shows only mild tail deviations.

Conclusion:

Although the Shapiro–Wilk test rejects normality, this is expected with large samples.

The residual distribution is *approximately normal*, with symmetry and unimodality clearly visible.

OLS estimates remain unbiased and efficient, and inference remains valid due to the Central Limit Theorem.

Overall Assessment

| Assumption | Test Result | Visual Assessment | Practical Verdict |
|------------------|----------------|-----------------------|-------------------|
| Independence | DW significant | No clear pattern | Acceptable |
| Homoscedasticity | BP significant | Mild variance changes | Acceptable |
| Normality | SW significant | Roughly symmetric | Acceptable |

Session Info

| Package | Version |
|-------------|---------|
| broom | 1.0.10 |
| car | 3.1-3 |
| collapse | 2.1.4 |
| correlation | 0.8.8 |
| datawizard | 1.3.0 |
| effectsize | 1.0.1 |
| GGally | 2.4.0 |
| ggfortify | 0.4.19 |
| ggpubr | 0.6.2 |
| glue | 1.8.0 |
| gtsummary | 2.4.0 |
| insight | 1.4.2 |
| kableExtra | 1.4.0 |
| lmtest | 0.9-40 |
| matrixTests | 0.2.3.1 |
| modelbased | 0.13.0 |
| parameters | 0.28.2 |
| patchwork | 1.3.2 |
| performance | 0.15.2 |
| qqplotr | 0.0.7 |
| rstatix | 0.7.3 |
| scales | 1.4.0 |
| see | 0.12.0 |
| tidyverse | 2.0.0 |